

Clarifying Question Dataset

Fine-tuning data for judging whether a developer prompt is executable and asking the single best clarifying question when it isn't.

Total examples	420
Executable	210
Needs clarification	210
missing_context	54
missing_requirement	26
missing_scope	25
missing_technology	16
multiple_interpretations	16
unclear_goal	15
missing_component	15
missing_acceptance_criteria	15
missing_file	14
missing_constraint	14

#1 EXECUTABLE

confidence 0.97

Add a --dry-run flag to the deploy script in scripts/deploy.sh that prints the commands it would run without executing them

#2 NEEDS CLARIFICATION type: missing_scope

confidence 0.96

make it faster

Ambiguities

- No target identified for optimization
- No performance baseline or goal stated

Missing information

- Which part of the system is slow
- What metric defines fast enough

Clarifying question: What specifically feels slow — page load, an API endpoint, a database query, or the build/test cycle?

#3 EXECUTABLE

confidence 0.95

Convert src/utils/dates.js from moment.js to date-fns and update the imports everywhere it's used

#4 NEEDS CLARIFICATION type: missing_context

confidence 0.97

fix the bug

Ambiguities

- No bug described
- No error message, repro steps, or affected feature given

Missing information

- What incorrect behavior is observed
- How to reproduce it

Clarifying question: What's the incorrect behavior you're seeing, and is there an error message or a way to reproduce it?

#5 EXECUTABLE

confidence 0.93

The tests in tests/test_auth.py started failing after my last commit. Run them, figure out what broke, and fix it.

#6 NEEDS CLARIFICATION type: missing_file

confidence 0.90

can you clean this up

Ambiguities

- No file or code selection referenced
- Cleanup criteria undefined

Missing information

- Which file or code region to clean up

Clarifying question: Which file or function do you want cleaned up?

#7 EXECUTABLE

confidence 0.96

Write a Dockerfile for this project. It's a Node 20 Express app, entry point is server.js, listens on port 3000.

#8 NEEDS CLARIFICATION type: missing_requirement

confidence 0.88

add auth to the app

Ambiguities

- Authentication method unspecified
- User store unspecified
- Unclear if third-party login is needed

Missing information

- Desired auth mechanism (sessions, JWT, OAuth provider, magic links)

Clarifying question: Do you want email/password auth with sessions, JWT-based auth for the API, or social login through a provider like Google?

#9 EXECUTABLE

confidence 0.94

I'm getting `TypeError: Cannot read properties of undefined (reading 'map')` in src/components/OrderList.tsx when the page first loads. Fix it.

#10 NEEDS CLARIFICATION type: missing_technology

confidence 0.92

deploy it

Ambiguities

- Deployment target unknown
- No deployment config found in conversation
- Unclear which environment

Missing information

- Hosting platform and environment to deploy to

Clarifying question: Where should this be deployed — do you have an existing setup like Vercel, AWS, or a VPS, and is this going to staging or production?

#11 EXECUTABLE

confidence 0.97

Rename the function getUserData to fetchUserProfile across the whole repo, including tests and any docs that mention it.

#12 **NEEDS CLARIFICATION** type: missing_context

confidence 0.85

the page is broken on mobile

Ambiguities

- Which page is unspecified
- Broken behavior undefined (layout, crash, interaction?)
- Device/browser unknown

Missing information

- Which page and what the broken behavior looks like

Clarifying question: Which page is broken, and what do you see — overlapping layout, missing content, or something not responding to taps?

#13 **EXECUTABLE**

confidence 0.96

Add an index on the created_at column of the orders table and write the migration using our existing Alembic setup in migrations/

#14 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.93

update the api

Ambiguities

- No endpoint specified
- Nature of the update unknown

Missing information

- Which endpoint(s) to change and what the change is

Clarifying question: Which endpoint needs updating, and what should change about it — the response shape, validation, behavior, or something else?

#15 **EXECUTABLE**

confidence 0.95

Set up ESLint and Prettier in this repo with the recommended TypeScript rules, and add a lint script to package.json

#16 **NEEDS CLARIFICATION** type: missing_context

confidence 0.80

why is prod slow today

Ambiguities

- No access to production metrics or logs from here
- Slowness not characterized (which requests, since when)

Missing information

- Observed symptoms: affected endpoints, latency numbers, or recent changes

Clarifying question: Which requests are slow and do you have any metrics or logs I can look at — e.g., p95 latency by endpoint or a recent deploy that correlates with the slowdown?

#17 **EXECUTABLE**

confidence 0.98

Write a Python script that reads orders.csv, groups rows by customer_id, sums the amount column per customer, and writes the result to totals.json

#18 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.94

delete old records from the events table

Ambiguities

- Definition of 'old' unspecified
- Destructive operation without a stated retention rule
- Unclear if archival is needed before deletion

Missing information

- Retention cutoff (age or date) for deletion

Clarifying question: What's the cutoff for 'old' — for example, delete events older than 90 days, and should they be archived anywhere first?

#19 **EXECUTABLE**

confidence 0.98

Add a healthcheck endpoint GET /healthz to the FastAPI app in app/main.py that returns 200 with {"status": "ok"}

#20 **NEEDS CLARIFICATION** type: missing_component

confidence 0.82

make the button blue

Ambiguities

- Which button is unspecified
- Exact shade of blue undefined

Missing information

- Which button/component in which view

Clarifying question: Which button do you mean — can you tell me the page or component it's on?

#21 **EXECUTABLE**

confidence 0.94

Refactor the duplicated retry logic in src/api/client.ts and src/api/uploads.ts into a shared helper, keeping the existing backoff behavior identical.

#22 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.95

migrate the database

Ambiguities

- Could mean run pending schema migrations, move to a different engine, or move to a new host
- Source and target not stated

Missing information

- Whether this is a schema migration, engine change, or host move

Clarifying question: Do you mean run the pending schema migrations, or migrate the data to a different database engine or host?

#23 **EXECUTABLE**

confidence 0.96

Our checkout API at POST /api/checkout returns 500 when the cart contains a discounted item. Stack trace points to apply_discounts() in services/pricing.py. Reproduce with the test fixture in tests/fixtures/discounted_cart.json and fix it.

#24 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.90

add caching

Ambiguities

- Caching layer unspecified (HTTP, application, database, CDN)
- What data to cache and for how long is unknown

Missing information

- Which operations or data should be cached

Clarifying question: What do you want cached — API responses, database query results, or static assets — and roughly how fresh does the data need to stay?

#25 **EXECUTABLE**

confidence 0.97

Write unit tests for the `parseConfig` function in `src/config/parser.ts` covering valid configs, missing required fields, and malformed YAML. We use Vitest.

#26 **NEEDS CLARIFICATION** type: missing_context

confidence 0.97

it doesn't work anymore

Ambiguities

- 'It' has no referent
- No description of failure or when it started

Missing information

- What stopped working and what changed before it broke

Clarifying question: What stopped working — which feature or command — and what's the error or behavior you see now?

#27 **EXECUTABLE**

confidence 0.96

Upgrade `lodash` to the latest 4.x in `package.json` and run the test suite to confirm nothing breaks.

#28 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.72

handle errors better in the worker

Ambiguities

- 'Better' is undefined
- Desired behavior on failure unknown (retry, dead-letter, alert, skip)

Missing information

- What should happen when a job fails

Clarifying question: When a worker job fails, what should happen — retry with backoff, send it to a dead-letter queue, alert someone, or skip and log?

#29 **EXECUTABLE**

confidence 0.97

Add pagination to `GET /api/products` using `limit/offset` query params, defaulting to `limit=20`, and include `total_count` in the response.

#30 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.85

integrate stripe

Ambiguities

- Payment model unspecified (one-time, subscriptions, marketplace)
- Checkout approach unspecified (Stripe Checkout vs custom Elements)

Missing information

- What is being sold and the desired payment flow

Clarifying question: Are you charging one-time payments or subscriptions, and do you want Stripe's hosted Checkout page or an embedded payment form?

#31 **EXECUTABLE**

confidence 0.97

Write a GitHub Actions workflow that runs pytest on every PR against main, using Python 3.12 and caching pip dependencies.

#32 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.70

set up ci

Ambiguities

- What CI should run is unspecified (tests, lint, build, deploy)
- CI provider not stated, though repo host may imply one

Missing information

- What checks the pipeline should run

Clarifying question: What should CI run on each push — just the test suite, or also linting, type checks, and a build step?

#33 **EXECUTABLE**

confidence 0.97

Our k8s deployment in k8s/api-deployment.yaml keeps getting OOMKilled. Bump the memory limit to 1Gi and the request to 512Mi.

#34 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.92

move us to kubernetes

Ambiguities

- Which services to migrate is unspecified
- Target cluster/platform unknown (EKS, GKE, self-hosted)
- Timeline and rollout strategy undefined

Missing information

- Which services move first and where the cluster will run

Clarifying question: Which service do you want to containerize and migrate first, and do you already have a cluster (EKS/GKE/AKS) or does that need provisioning too?

#35 **EXECUTABLE**

confidence 0.95

Squash the last 4 commits on this branch into one with the message 'Add rate limiting to public API'

#36 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.85

undo my changes

Ambiguities

- Could mean discard uncommitted changes, revert a commit, or reset the branch
- Scope unknown (one file, all files, last commit)

Missing information

- Which changes to undo and whether they are committed

Clarifying question: Do you want to discard your uncommitted edits, or undo commits that are already made — and if commits, how far back?

#37 **EXECUTABLE**

confidence 0.98

Add a .gitignore appropriate for a Python project using venv, plus entries for .env and the data/ directory.

#38 **NEEDS CLARIFICATION** type: missing_context

confidence 0.60

the docker build is broken

Ambiguities

- Error output not provided, though the build may be reproducible locally
- Which image/Dockerfile if there are several

Missing information

- The build error output or which Dockerfile fails

Clarifying question: Can I reproduce it by running `docker build .` here, or does it only fail in CI — and if so, can you paste the error from the CI log?

#39 **EXECUTABLE**

confidence 0.97

Write a Terraform module for an S3 bucket with versioning enabled, server-side encryption with SSE-S3, and public access fully blocked.

#40 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

we need monitoring

Ambiguities

- What to monitor is unspecified (uptime, errors, latency, infra metrics)
- Tooling preference unknown
- Alerting requirements undefined

Missing information

- What signals matter most and any existing observability stack

Clarifying question: What's the most important thing to catch first — service downtime, error spikes, or slow requests — and do you already use any tool like Datadog, Grafana, or Sentry?

#41 **EXECUTABLE**

confidence 0.93

Cherry-pick commit 3f2a91c from main onto the release/2.4 branch and resolve any conflicts in favor of the release branch's config files.

#42 **NEEDS CLARIFICATION** type: missing_context

confidence 0.65

can you make the pipeline not flaky

Ambiguities

- Which jobs/tests are flaky is unspecified, though CI history may reveal it
- Flakiness pattern unknown (timeouts, race conditions, infra)

Missing information

- Which jobs or tests fail intermittently

Clarifying question: Which job or test is the flaky one — or should I pull the last 20 CI runs and identify the most frequent intermittent failures myself?

#43 **EXECUTABLE**

confidence 0.97

Add a pre-commit hook that runs ruff and black on staged Python files, using the pre-commit framework.

#44 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.90

rotate the secrets

Ambiguities

- Which secrets to rotate is unspecified
- Where secrets live is unknown (env files, vault, cloud secret manager)
- Rotation may break running services if done without coordination

Missing information

- Which credentials to rotate and where they are stored/consumed

Clarifying question: Which secrets need rotating — a specific leaked key, or everything — and are they stored in a secrets manager I can access?

#45 **EXECUTABLE**

confidence 0.95

Our nginx config at /etc/nginx/sites-available/api proxies to localhost:8000. Add gzip compression for JSON responses and a 30s proxy read timeout.

#46 **NEEDS CLARIFICATION** type: missing_context

confidence 0.68

make deploys faster

Ambiguities

- Current deploy process and duration unknown, though CI config may show it
- Which stage is slow is unidentified

Missing information

- What the deploy pipeline looks like and where the time goes

Clarifying question: Roughly how long does a deploy take now, and do you know which stage eats most of it — build, tests, image push, or the rollout itself?

#47 **EXECUTABLE**

confidence 0.97

Create a docker-compose.yml for local dev with the api service built from ./Dockerfile, a postgres:16 container with a named volume, and the API getting DATABASE_URL pointed at it.

#48 **NEEDS CLARIFICATION** type: missing_context

confidence 0.91

set up staging

Ambiguities

- Current hosting setup unknown
- Whether staging should mirror prod infra or be lightweight is undefined
- Data strategy for staging unspecified

Missing information

- Where production runs and what staging needs to replicate

Clarifying question: Where does production run today, and should staging be a full mirror of it or a cheaper scaled-down environment?

#49 **EXECUTABLE**

confidence 0.96

Write an Ansible playbook that installs nginx on Ubuntu 22.04 hosts in the webserver group and copies templates/nginx.conf.j2 to /etc/nginx/nginx.conf

#50 **NEEDS CLARIFICATION** type: missing_context

confidence 0.62

the cron job didn't run last night

Ambiguities

- Which cron job and where it runs is unspecified
- Logs location unknown, though system crontab may be inspectable

Missing information

- Which job, and where it's scheduled (system cron, k8s CronJob, scheduler service)

Clarifying question: Which job is it and where is it scheduled — a system crontab on a server, a Kubernetes CronJob, or something like Celery beat?

#51 **EXECUTABLE**

confidence 0.95

Pin all GitHub Actions in .github/workflows/ to full commit SHAs instead of version tags.

#52 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

reduce our aws bill

Ambiguities

- No visibility into current spend breakdown
- Acceptable tradeoffs unknown (performance, redundancy)
- Which services dominate cost is unidentified

Missing information

- Cost breakdown by service and what tradeoffs are acceptable

Clarifying question: Do you have a Cost Explorer breakdown showing which services dominate the bill, or can you give me read access so I can find the biggest line items first?

#53 **EXECUTABLE**

confidence 0.98

Add a Makefile with targets: install (pip install -r requirements.txt), test (pytest), lint (ruff check .), and format (black .).

#54 **NEEDS CLARIFICATION** type: missing_context

confidence 0.86

we should do blue green deployments

Ambiguities

- Current deployment infrastructure unknown
- Traffic switching mechanism depends on platform (LB, ingress, service mesh)

Missing information

- What platform serves traffic today and how deploys currently work

Clarifying question: What does your current deploy target look like — ECS, Kubernetes, VMs behind a load balancer — so I know what would do the traffic switching?

#55 **EXECUTABLE**

confidence 0.92

Tag the current commit as v2.3.0, push the tag, and draft release notes from the commits since v2.2.0.

#56 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.89

lock down the server

Ambiguities

- Which server is unspecified
- Threat model and hardening level undefined
- Risk of locking out legitimate access

Missing information

- Which host and what prompted the hardening

Clarifying question: Which server, and is this routine hardening or a response to something specific — that changes whether I start with SSH/firewall basics or hunt for an active issue?

#57 **EXECUTABLE**

confidence 0.95

In .github/workflows/ci.yml, split the single job into two parallel jobs: one for lint+typecheck, one for tests, both required for merge.

#58 **NEEDS CLARIFICATION** type: missing_context

confidence 0.75

i accidentally pushed a secret to the repo

Ambiguities

- Which secret and which commits is unspecified, though git history can be searched
- Whether the repo is public/shared affects urgency
- Rotation owner unknown

Missing information

- What kind of secret it is and whether the repo is public

Clarifying question: What kind of secret is it (API key, DB password) and is the repo public — rotating the credential comes first, so what does it unlock?

#59 **EXECUTABLE**

confidence 0.96

Update the Helm chart in charts/api to add a PodDisruptionBudget with minAvailable: 1 for the api deployment.

#60 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.80

automate the releases

Ambiguities

- What 'release' means here is unclear (npm publish, GitHub release, deploy, app store)
- Versioning policy unspecified
- Trigger condition unknown (tag, merge to main, manual)

Missing information

- What artifact gets released and what should trigger it

Clarifying question: What does a release produce for you — an npm package, a GitHub release with binaries, a deployed service — and should it trigger on every merge to main or on tags?

#61 **EXECUTABLE**

confidence 0.96

Add a loading spinner to the SubmitButton component in src/components/SubmitButton.tsx while the form's onSubmit promise is pending.

#62 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.95

redesign the dashboard

Ambiguities

- No design direction given
- Unclear what's wrong with the current dashboard
- Scope could range from styling tweaks to full information architecture changes

Missing information

- What problems the redesign should solve or what it should look like

Clarifying question: What's driving the redesign — is the current dashboard cluttered, missing key data, visually dated — and do you have a reference design or mockup?

#63 **EXECUTABLE**

confidence 0.92

The navbar dropdown in src/components/Nav.vue closes immediately when you click an item on Safari but works in Chrome. Likely a blur/mousedown ordering issue — fix it.

#64 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.95

make it look better

Ambiguities

- 'Better' is subjective with no direction given
- Target page/component unspecified

Missing information

- Which screen and what aesthetic direction (examples, brand guidelines)

Clarifying question: Which screen are we improving, and is there a site or design you'd want it to feel more like?

#65 **EXECUTABLE**

confidence 0.95

Convert the class component in src/components/UserSettings.jsx to a function component with hooks, preserving all existing behavior including the componentDidUpdate logic.

#66 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.60

add dark mode

Ambiguities

- Toggle behavior unspecified (manual toggle, follow OS, both)
- Persistence approach undefined
- Whether a design system/tokens already exist affects approach

Missing information

- Whether dark mode should follow the OS setting or be a manual toggle

Clarifying question: Should dark mode follow the user's OS preference automatically, be a manual toggle, or both with the toggle overriding?

#67 **EXECUTABLE**

confidence 0.97

Center the login card vertically and horizontally on the page in `src/pages/Login.tsx` – it currently sits at the top-left. We use Tailwind.

#68 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

the form is janky

Ambiguities

- Which form is unspecified
- 'Janky' could mean laggy input, layout shift, or odd validation behavior

Missing information

- Which form and what the jankiness looks like

Clarifying question: Which form, and what does the jank look like — typing lag, the layout jumping around, or validation firing at weird times?

#69 **EXECUTABLE**

confidence 0.97

Add client-side validation to the signup form in `src/pages/Signup.tsx`: email format, password min 8 chars with at least one number, and show inline error messages below each field.

#70 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.62

switch to tailwind

Ambiguities

- Whether to convert all existing styles or just adopt for new code
- Fate of existing CSS/styled-components unclear
- Big-bang vs incremental migration undefined

Missing information

- Migration scope: full conversion vs incremental adoption

Clarifying question: Do you want the entire existing codebase converted to Tailwind now, or just Tailwind set up so new components use it while old styles migrate gradually?

#71 **EXECUTABLE**

confidence 0.94

Our Next.js 14 app router page at `app/products/[id]/page.tsx` fetches client-side. Convert it to a server component with the fetch happening on the server, keeping the AddToCart button as a client component.

#72 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.55

improve accessibility

Ambiguities

- Whole app vs specific flows
- Target standard unstated (WCAG 2.1 AA is a reasonable default)
- Whether this is audit-driven or proactive

Missing information

- Which pages/flows to prioritize and any compliance target

Clarifying question: Is there a specific compliance bar you need to hit (like WCAG 2.1 AA for a customer contract), or should I audit the main user flows and fix the worst issues first?

#73 **EXECUTABLE**

confidence 0.97

Debounce the search input in `src/components/SearchBar.tsx` by 300ms so we stop firing an API call on every keystroke.

#74 **NEEDS CLARIFICATION** type: missing_context

confidence 0.58

users say the site feels slow

Ambiguities

- Subjective report without specifics, though Lighthouse/profiling can be run
- Which pages/flows users mean is unknown
- Could be network, rendering, or backend latency

Missing information

- Which interactions feel slow and on what kind of device/connection

Clarifying question: Did users mention specific pages or actions that feel slow, or should I start by profiling the highest-traffic pages with Lighthouse and work from the worst scores?

#75 **EXECUTABLE**

confidence 0.95

Replace the moment-based relative timestamps in `src/components/CommentList.tsx` with `Intl.RelativeTimeFormat` and delete the moment dependency if nothing else uses it.

#76 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.90

add a chart to the analytics page

Ambiguities

- What data the chart should show is unspecified
- Chart type undefined
- Data source/endpoint unknown

Missing information

- What metric to visualize and over what time range

Clarifying question: What should the chart show — which metric, over what time range, and does an API endpoint for that data already exist?

#77 **EXECUTABLE**

confidence 0.95

Extract the inline styles in `src/pages/Pricing.jsx` into a CSS module, keeping the rendered output pixel-identical.

#78 **NEEDS CLARIFICATION** type: missing_context

confidence 0.93

the modal is doing the thing again

Ambiguities

- 'The thing' refers to prior context not available
- Which modal is unspecified

Missing information

- What behavior 'the thing' refers to and which modal

Clarifying question: Which modal, and what is it doing — I don't have the earlier context, so can you describe the bad behavior you're seeing?

#79 **EXECUTABLE**

confidence 0.97

Add keyboard navigation to the autocomplete in src/components/Autocomplete.tsx: arrow keys to move through results, Enter to select, Escape to close.

#80 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.92

make the app work offline

Ambiguities

- Offline scope undefined (read-only cache vs full offline editing with sync)
- Which features must work offline is unspecified
- Conflict resolution for offline writes is a major design decision

Missing information

- Which functionality must work offline and whether offline writes are needed

Clarifying question: Does offline mean users can view previously loaded content, or do they need to create and edit data offline that syncs back later?

#81 **EXECUTABLE**

confidence 0.96

Lazy-load the routes in src/App.tsx with React.lazy and Suspense, using the existing PageSpinner component as the fallback.

#82 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.70

i want it to look like linear

Ambiguities

- Which aspects of Linear's design to emulate (typography, density, dark theme, animations)
- Which screens to restyle first

Missing information

- Which elements of Linear's aesthetic matter most and which screens to start with

Clarifying question: What is it about Linear you want most — the dense compact layout, the dark theme and typography, or the snappy keyboard-driven feel — and which screen should I restyle first as the reference?

#83 **EXECUTABLE**

confidence 0.93

Fix the hydration mismatch warning in app/layout.tsx — it's caused by rendering new Date() directly during SSR.

#84 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.82

add i18n

Ambiguities

- Target languages unspecified
- Whether translations exist or need placeholder extraction only
- Locale detection/switching behavior undefined

Missing information

- Which languages to support initially

Clarifying question: Which languages do you need at launch, and do you already have translations or should I just externalize the strings so a translator can fill them in?

#85 **EXECUTABLE**

confidence 0.93

Build a 404 page matching our existing design system components in `src/components/ui/`, with a link back to the home page.

#86 **NEEDS CLARIFICATION** type: missing_component

confidence 0.85

the dropdown should work like the other one

Ambiguities

- Neither dropdown is identified
- Which behaviors to copy is unspecified

Missing information

- Which dropdown to change and which one is the reference

Clarifying question: Which dropdown needs fixing, and which existing one is the reference for how it should behave?

#87 **EXECUTABLE**

confidence 0.96

Memoize the expensive `filteredProducts` computation in `src/pages/Catalog.tsx` with `useMemo` – it currently recomputes on every render and the list has ~5k items.

#88 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.96

we need the landing page done by friday

Ambiguities

- No content, copy, or design provided
- Sections and structure undefined
- Unclear what 'done' includes (responsive, SEO, analytics)

Missing information

- Copy/design assets and required sections

Clarifying question: Do you have copy and a design for it, or should I draft the structure — and at minimum, what sections must it have (hero, pricing, testimonials, signup form)?

#89 **EXECUTABLE**

confidence 0.96

Stop the layout shift on the article pages: the hero image in `src/components/ArticleHero.astro` has no dimensions, so text jumps when it loads. Add width/height and aspect-ratio CSS.

#90 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.66

hook up the frontend to the backend

Ambiguities

- Which features/pages to wire up is unspecified
- API contract may or may not exist
- Auth handling between the two is undefined

Missing information

- Which screens need live data first and whether the API endpoints for them exist

Clarifying question: Which screen should I wire up first, and do the API endpoints for it already exist or do I need to build those too?

#91 **EXECUTABLE**

confidence 0.97

Add rate limiting to the POST /api/login endpoint in our Express app: max 5 attempts per IP per 15 minutes, return 429 with a Retry-After header.

#92 **NEEDS CLARIFICATION** type: missing_component

confidence 0.80

the query is slow

Ambiguities

- Which query is unspecified
- No EXPLAIN output or timing provided
- Table sizes and existing indexes unknown

Missing information

- The query text or where it lives in code

Clarifying question: Which query — can you point me to where it's defined in code, or paste it so I can run EXPLAIN ANALYZE on it?

#93 **EXECUTABLE**

confidence 0.94

Write a SQL migration that adds a nullable phone_number varchar(20) column to the users table, then a follow-up that backfills it from the legacy contacts table joined on user_id.

#94 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.85

normalize the database

Ambiguities

- Which tables have normalization problems is unspecified, though schema can be inspected
- Degree of normalization desired vs read-performance tradeoffs
- Whether app code changes are in scope

Missing information

- Which denormalized data is causing actual problems

Clarifying question: Is there a specific table where duplicated data is causing bugs or update anomalies, or do you want me to review the whole schema and propose changes before touching anything?

#95 **EXECUTABLE**

confidence 0.95

Add a unique constraint on (user_id, course_id) in the enrollments table. First write a query to find existing duplicates so we can decide how to handle them.

#96 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.91

dedupe the users table

Ambiguities

- What makes two users duplicates is undefined (same email? name? fuzzy match?)
- Which record survives a merge is unspecified
- Foreign keys referencing removed rows need a strategy

Missing information

- The duplicate-matching rule and which record to keep

Clarifying question: What counts as a duplicate — exact same email, or fuzzier matching — and when two match, should I keep the oldest account and repoint the other's data to it?

#97 **EXECUTABLE**

confidence 0.96

Our GET /api/orders endpoint does an N+1: it fetches orders then queries line items per order in a loop in services/orders.py. Fix it with a joined or batched query.

#98 **NEEDS CLARIFICATION** type: missing_context

confidence 0.70

should we use mongo or postgres

Ambiguities

- Data model and access patterns unknown
- Scale and consistency requirements unstated
- Team experience not mentioned

Missing information

- What kind of data and queries the app needs

Clarifying question: What does your core data look like — highly relational entities with joins and transactions, or flexible documents queried mostly by ID?

#99 **EXECUTABLE**

confidence 0.96

Add a soft-delete to the Post model: a deleted_at timestamp, default scope excluding deleted rows, and a .deleted scope to find them. It's a Rails 7 app.

#100 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.94

build the api for the mobile app

Ambiguities

- No feature list or endpoints specified
- Data model unknown
- Auth approach for mobile clients undefined

Missing information

- What the mobile app does and what endpoints it needs

Clarifying question: What are the core screens of the mobile app, so I can derive the endpoints — or do you already have a list of the API operations it needs?

#101 **EXECUTABLE**

confidence 0.97

Change the /api/users/:id response to include the user's organization name. The Organization model is already associated via user.org_id — just join and add org_name to the serializer in serializers/user.py.

#102 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.78

version the api

Ambiguities

- Versioning scheme unspecified (URL path, header, query param)
- Whether existing clients must keep working unchanged
- Whether this is preemptive or driven by a planned breaking change

Missing information

- The versioning scheme preference and what breaking change is coming

Clarifying question: Is there a specific breaking change driving this, and do you prefer /v1/ URL prefixes or header-based versioning?

#103 **EXECUTABLE**

confidence 0.97

Write a database seed script in prisma/seed.ts that creates 3 users, 2 organizations, and 10 sample projects with realistic fake names using faker.

#104 **NEEDS CLARIFICATION** type: missing_context

confidence 0.62

the db connections keep maxing out

Ambiguities

- Could be pool misconfiguration, connection leaks, or genuine load — code inspection can narrow it
- Database and pool settings unknown
- When it happens (steady state vs spikes) unstated

Missing information

- When the exhaustion happens and current pool configuration

Clarifying question: Does it happen under load spikes or build up gradually over hours — gradual buildup points to a connection leak, which I can hunt for in the code?

#105 **EXECUTABLE**

confidence 0.95

Add cursor-based pagination to the GraphQL posts query in src/graphql/resolvers/posts.ts following the Relay connection spec we already use for comments.

#106 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.68

make the api restful

Ambiguities

- Which conventions are being violated is unspecified, though routes can be audited
- Breaking existing clients may not be acceptable
- Depth of restructuring undefined

Missing information

- Whether breaking changes to existing consumers are acceptable

Clarifying question: Can I break the existing routes, or do current clients need to keep working — that decides between restructuring the API and adding a parallel set of conventional routes?

#107 **EXECUTABLE**

confidence 0.94

Replace the raw SQL in repositories/user_repo.go with sqlc-generated queries. The queries live in db/queries/users.sql and sqlc.yaml is already configured.

#108 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

add webhooks

Ambiguities

- Direction unclear: consuming external webhooks vs letting our users subscribe to ours
- Which events to emit/handle is unspecified

Missing information

- Whether we send or receive webhooks, and for which events

Clarifying question: Do you mean receiving webhooks from a third party like Stripe, or building outgoing webhooks so your customers can subscribe to your events — and which events?

#109 **EXECUTABLE**

confidence 0.93

Set up connection pooling for our Django app with pgbouncer in transaction mode; update settings.py and add the pgbouncer service to docker-compose.yml.

#110 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.87

store the files somewhere

Ambiguities

- Storage target unspecified (S3, GCS, local disk, database)
- Which files and their size/access patterns unknown
- Public vs private access undefined

Missing information

- What files these are and where infrastructure already exists

Clarifying question: What kind of files and roughly how large — and are you already on a cloud provider whose storage (like S3) we should use?

#111 **EXECUTABLE**

confidence 0.96

Write an idempotent upsert for the daily_metrics table keyed on (date, metric_name) — Postgres INSERT ... ON CONFLICT DO UPDATE — and use it in jobs/aggregate_metrics.py.

#112 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.83

the api should be async

Ambiguities

- Could mean async/await internals, background job processing, or async request/callback API design
- Motivation unknown (throughput? long-running tasks?)

Missing information

- Which sense of 'async' is intended and what problem it solves

Clarifying question: Do you mean converting the handlers to async/await for throughput, or making slow operations return immediately and finish in a background job the client polls?

#113 **EXECUTABLE**

confidence 0.97

Add request validation to POST /api/projects using zod: name required 1-100 chars, description optional max 2000, tags an array of max 10 strings. Return 422 with field-level errors.

#114 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.90

we're switching to graphql

Ambiguities

- Full replacement vs GraphQL alongside REST
- Which clients migrate and when
- Schema design ownership unclear

Missing information

- Migration scope and which API surface goes first

Clarifying question: Is GraphQL replacing the REST API entirely, or running alongside it — and which feature area should get the first schema?

#115 **EXECUTABLE**

confidence 0.94

Profile why POST /api/import takes 30s for a 10MB CSV — add timing logs around the parse, validation, and insert phases in handlers/import.go, run the sample file in testdata/, and report where the time goes.

#116 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.71

cache the user data

Ambiguities

- Staleness tolerance unspecified
- Invalidation strategy on user updates undefined
- Cache store (in-memory vs Redis) depends on deployment shape

Missing information

- How stale user data can be and what must invalidate it

Clarifying question: How stale can cached user data be — is a 5-minute TTL fine, or must changes like a role update take effect immediately, which means cache invalidation on write?

#117 **EXECUTABLE**

confidence 0.95

Convert the Express error handling to a central middleware: catch async errors, log with our pino logger, return JSON {error, requestId}, and map known error classes in src/errors.ts to status codes.

#118 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.94

make sure we don't lose data

Ambiguities

- Threat unspecified (crashes, bad deploys, accidental deletion, disk failure)
- Could mean backups, transactions, soft deletes, or replication
- Current data infrastructure unknown

Missing information

- What data-loss scenario is the concern

Clarifying question: What scenario worries you — accidental deletes by users, a bad migration wiping rows, or losing the database server itself? Each needs a different safeguard.

#119 **EXECUTABLE**

confidence 0.95

Add a GET /api/search endpoint that does full-text search over posts.title and posts.body using Postgres tsvector, ranked by ts_rank, limited to 25 results.

#120 **NEEDS CLARIFICATION** type: missing_component

confidence 0.65

the orm is generating bad queries

Ambiguities

- Which queries/endpoints are affected is unspecified, though query logging could find candidates
- 'Bad' undefined (slow, wrong results, too many round trips)

Missing information

- Which operation produces the bad query and what's bad about it

Clarifying question: Which endpoint or operation produces the bad query, and is the problem slowness, N+1 round trips, or incorrect results?

#121 **EXECUTABLE**

confidence 0.90

Run the test suite and fix whatever's failing.

#122 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.60

write tests

Ambiguities

- No target module specified, though coverage gaps are discoverable
- Test type unspecified (unit, integration, e2e)
- Repo may have multiple untested areas with different priority

Missing information

- Which code to test first

Clarifying question: Is there a specific module you're worried about, or should I check coverage and write tests for the least-covered critical paths?

#123 **EXECUTABLE**

confidence 0.96

Add a regression test for the bug we just fixed: submitting the checkout form with an expired coupon should return 400, not apply a zero discount. Put it in tests/api/test_checkout.py.

#124 **NEEDS CLARIFICATION** type: missing_context

confidence 0.70

why does this fail in ci but pass locally

Ambiguities

- Which test/job is failing is unspecified, though CI logs may be accessible
- Environment differences unknown (versions, timezone, parallelism)

Missing information

- Which test fails and the CI failure output

Clarifying question: Which test is failing — and can you paste the CI failure output or give me the link to the failed run so I can compare it to a local run?

#125 **EXECUTABLE**

confidence 0.95

Mock the Stripe client in tests/test_billing.py so tests stop hitting the real API — use pytest-mock and assert the charge call gets the right amount and currency.

#126 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.55

get coverage to 80%

Ambiguities

- Coverage metric unspecified (line vs branch)
- Whether to exclude generated/vendored code
- Whether shallow tests just to hit the number are acceptable

Missing information

- What code counts toward the target and quality expectations for the new tests

Clarifying question: Is 80% line coverage across everything, or should generated code and migrations be excluded — and do you care that tests are meaningful rather than just executing lines?

#127 **EXECUTABLE**

confidence 0.95

Add Playwright e2e coverage for the happy-path signup flow: visit /signup, fill the form, submit, assert redirect to /onboarding and the welcome toast. The dev server runs via npm run dev.

#128 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

this is throwing somewhere

Ambiguities

- 'This' has no referent
- No stack trace or error text provided
- When it throws is unknown

Missing information

- The error message/stack trace and what action triggers it

Clarifying question: What's the error message or stack trace, and what were you doing when it happened?

#129 **EXECUTABLE**

confidence 0.90

The test test_concurrent_checkout in tests/test_orders.py is flaky – fails about 1 in 10 runs with a deadlock. Run it in a loop to reproduce, then fix the underlying race.

#130 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.62

review my code

Ambiguities

- Which code is unspecified, though uncommitted/branch diff is a reasonable default
- Review focus undefined (bugs, style, security, performance)

Missing information

- Whether to review the current branch diff or something else

Clarifying question: Should I review the diff on your current branch against main, or a different set of changes?

#131 **EXECUTABLE**

confidence 0.93

Bisect to find which commit between v1.8.0 and HEAD broke the import command – `python -m app import sample.csv` should exit 0. Use git bisect with that as the test.

#132 **NEEDS CLARIFICATION** type: missing_context

confidence 0.68

users are seeing 500s sometimes

Ambiguities

- Which endpoints and how often is unspecified
- No error logs or tracebacks provided
- 'Sometimes' pattern unknown (load-related, specific inputs)

Missing information

- Error logs or which endpoints return 500s

Clarifying question: Do you have error logs or a tracker like Sentry I can check — or at least which endpoints are returning the 500s?

#133 **EXECUTABLE**

confidence 0.95

Refactor tests/utils.py: the setup helpers create_test_user and create_test_org are duplicated in three test files. Consolidate them into conftest.py fixtures.

#134 **NEEDS CLARIFICATION** type: missing_file

confidence 0.75

is this code good

Ambiguities

- No code attached or file referenced
- Evaluation criteria undefined

Missing information

- Which code to evaluate

Clarifying question: Which file or snippet should I look at — and is there a specific concern, like readability, performance, or correctness?

#135 **EXECUTABLE**

confidence 0.96

Add property-based tests with Hypothesis for the slugify function in utils/text.py: output is always lowercase, URL-safe, and idempotent.

#136 **NEEDS CLARIFICATION** type: missing_context

confidence 0.60

something broke after the merge

Ambiguities

- What 'broke' means is unspecified, though tests and the merge diff can be examined
- Which merge if several recent ones

Missing information

- The observed breakage symptom

Clarifying question: What's the symptom — failing tests, a build error, or wrong behavior at runtime? If it's tests or build, I can run them and diff against the merge.

#137 **EXECUTABLE**

confidence 0.93

Snapshot tests in src/components/__tests__ keep churning on every PR because of generated IDs. Replace the random ID in src/utils/id.ts with a seeded/mockable generator in tests.

#138 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.72

test it like a real user would

Ambiguities

- Which flows constitute real user behavior is unspecified
- Tooling preference unknown (Playwright, Cypress)
- Environments/browsers to cover undefined

Missing information

- Which user journeys matter most

Clarifying question: Which two or three user journeys are most critical to protect — for example signup, checkout, and search — so I start the e2e suite there?

#139 **EXECUTABLE**

confidence 0.92

Find and delete dead code: run knip on this TypeScript repo, review what it flags, and remove exports nothing imports — but leave the public API in src/index.ts alone.

#140 **NEEDS CLARIFICATION** type: missing_context

confidence 0.66

the memory keeps growing

Ambiguities

- Which process/service leaks is unspecified
- Runtime unknown (Node, Python, JVM each profile differently)
- Growth rate and conditions unstated

Missing information

- Which service leaks and under what conditions

Clarifying question: Which service is it, and does memory grow steadily under normal traffic or only during specific operations like large imports?

#141 **EXECUTABLE**

confidence 0.92

Set up mutation testing with Stryker for the src/billing module and report the mutation score — don't fix anything yet, just give me the survivors.

#142 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.58

double check everything before we ship

Ambiguities

- 'Everything' is unbounded
- What shipping entails here is unknown
- Which risks matter most (regressions, security, performance) is unstated

Missing information

- What's shipping and which failure would be most costly

Clarifying question: What's in this release, and what would be the most painful thing to get wrong — so I prioritize checking that first rather than skimming everything?

#143 **EXECUTABLE**

confidence 0.94

Reproduce issue #482: per the report, calling the CLI with --config pointing to a symlinked file crashes with ENOENT. Write a failing test, then fix the path resolution in src/config/load.ts.

#144 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

why is this null

Ambiguities

- 'This' has no referent — no variable, file, or stack trace given
- Runtime state not observable without repro

Missing information

- Which value is null and where it's observed

Clarifying question: Which variable is null and where are you seeing it — a file and line, or a log statement I can look at?

#145 **EXECUTABLE**

confidence 0.94

Convert these Jest tests in `src/__tests__` to Vitest — the `vitest.config.ts` already exists; update imports, timer mocks, and the npm test script, then make sure they all pass.

#146 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.70

make the integration tests not need the real database

Ambiguities

- Could mean mocking the DB layer, an in-memory substitute, or ephemeral containers like testcontainers
- Fidelity vs speed tradeoff is a real decision
- Why the real DB is a problem is unstated (speed, CI access, data pollution)

Missing information

- What problem the real database causes and how much fidelity tests need

Clarifying question: What's the pain with the real DB — CI can't reach it, tests are slow, or data pollution? Spinning up a throwaway Postgres in Docker per test run keeps fidelity; mocking the data layer is faster but weaker.

#147 **EXECUTABLE**

confidence 0.94

Run ruff with `--fix` across the repo, then fix by hand whatever it can't auto-fix, and make sure the test suite still passes.

#148 **NEEDS CLARIFICATION** type: missing_context

confidence 0.63

it works on my machine but not for anyone else

Ambiguities

- What 'it' is and how it fails for others is unspecified
- Failure output from another machine not provided
- Difference candidates: env vars, versions, OS, uncommitted files

Missing information

- The error others see and how they're running it

Clarifying question: What error do the others get exactly — can you paste their output? Meanwhile I can check for uncommitted files and undocumented env vars your machine might have.

#149 **EXECUTABLE**

confidence 0.94

Add a `tests/load/k6_checkout.js` k6 script that ramps to 200 VUs over 2 minutes against staging's checkout flow and fails if p95 exceeds 800ms.

#150 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.68

stop the tests from being so slow

Ambiguities

- Acceptable tradeoffs unknown (parallelism, fewer integration tests, shared fixtures)
- Current duration and the slowest tests are discoverable but the target time isn't
- CI vs local speed may matter differently

Missing information

- Target duration and which tradeoffs are acceptable

Clarifying question: What would be fast enough — and are you open to tradeoffs like running tests in parallel or sharing database fixtures, which can introduce test-isolation risk?

#151 **EXECUTABLE**

confidence 0.97

Add retry with exponential backoff to the Claude API calls in `src/llm/client.py` — retry on 429 and 529 up to 5 times, starting at 1s with jitter.

#152 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.92

the model isn't good enough

Ambiguities

- Which model and which metric is unspecified
- 'Good enough' threshold undefined
- Failure examples not provided

Missing information

- Which metric is below target and example failures

Clarifying question: What does it get wrong — can you share a few failing examples and the metric you're judging it on, so I know whether to fix data, prompts, or the model choice?

#153 **EXECUTABLE**

confidence 0.97

Write a script to split `data/labeled.jsonl` into `train/val/test` at 80/10/10, stratified by the label field, with a fixed random seed, outputting three files in `data/splits/`.

#154 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.97

add ai to the product

Ambiguities

- No use case specified
- What user problem AI should solve is unknown
- Could be anything from a chatbot to search to recommendations

Missing information

- What user problem or feature the AI should power

Clarifying question: What should the AI actually do for users — answer questions about their data, summarize content, automate a workflow step, or something else you have in mind?

#155 **EXECUTABLE**

confidence 0.94

Our RAG pipeline in `src/rag/retrieve.py` fetches top-10 chunks by cosine similarity. Add a reranking step using the chunks' BM25 score blended 50/50 with the embedding score, and make `k` configurable.

#156 **NEEDS CLARIFICATION** type: missing_context

confidence 0.78

the chatbot hallucinates

Ambiguities

- No example transcripts provided
- Whether it invents facts, cites wrong sources, or answers outside its knowledge is unclear
- Acceptable behavior on unknown questions is undefined

Missing information

- Example bad transcripts and what it should say when it doesn't know

Clarifying question: Can you share two or three transcripts where it hallucinated — and when it genuinely doesn't know, should it say so, or escalate to a human?

#157 **EXECUTABLE**

confidence 0.95

Add token usage logging to every LLM call in src/agents/: log model, input/output tokens, latency, and calling function to our existing structlog setup, and write a daily cost rollup query.

#158 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

fine tune a model on our data

Ambiguities

- Task/objective unspecified (classification, generation, style)
- Which data and whether it's labeled is unknown
- Base model and budget unstated
- Fine-tuning may not even be the right tool vs prompting/RAG

Missing information

- What the fine-tuned model should do and what training data exists

Clarifying question: What should the model do after fine-tuning, and what data do you have — labeled examples of that exact task, or raw documents? If it's raw documents, RAG might fit better than fine-tuning.

#159 **EXECUTABLE**

confidence 0.94

Build an eval harness: run the prompts in evals/cases.yaml against our extract_entities() function, score exact-match per field, and print a table of per-field accuracy plus regressions vs the last run stored in evals/baseline.json.

#160 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.85

make the embeddings better

Ambiguities

- 'Better' undefined — retrieval relevance, clustering quality, or speed
- No eval set to measure improvement against
- Current model and chunking unknown without inspection

Missing information

- What downstream task the embeddings underperform on

Clarifying question: Where do the embeddings fall short — are users getting irrelevant search/RAG results, and do you have example queries where the wrong chunks come back?

#161 **EXECUTABLE**

confidence 0.96

Parallelize the embedding backfill in scripts/backfill_embeddings.py with a worker pool of 8, batch the API calls 100 texts at a time, and add a resume checkpoint so a crash doesn't restart from zero.

#162 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

should we build or buy for the llm stuff

Ambiguities

- What 'the llm stuff' is — which capability is being considered
- Constraints unknown (budget, data privacy, team size, latency)
- What 'buy' candidates are on the table

Missing information

- The specific capability and the constraints that matter

Clarifying question: Which capability is this decision about — and what's the binding constraint: cost, data not leaving your infra, or engineering time?

#163 **EXECUTABLE**

confidence 0.95

Convert the training loop in train.py to use mixed precision with torch.cuda.amp, and add gradient clipping at max_norm=1.0.

#164 **NEEDS CLARIFICATION** type: missing_context

confidence 0.62

the pipeline keeps breaking on weird data

Ambiguities

- Which pipeline stage fails is unspecified, though error logs may exist
- Example failing records not provided
- Desired handling of bad rows undefined (skip, quarantine, fail loudly)

Missing information

- Example failing inputs and the desired behavior for them

Clarifying question: Can you point me at a failing run or example bad records — and when a row is malformed, should the pipeline skip and log it, quarantine it for review, or halt?

#165 **EXECUTABLE**

confidence 0.95

Write a dbt model marts/fct_daily_signups.sql that counts signups per day per acquisition channel from stg_users, with a schema test that channel is never null.

#166 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

build a recommendation engine

Ambiguities

- What to recommend and where it surfaces is unspecified
- Available signals unknown (purchases, views, ratings)
- Scale and latency requirements unstated
- Cold-start expectations undefined

Missing information

- What gets recommended, from what interaction data

Clarifying question: What are we recommending and what interaction data exists — purchase history, page views, explicit ratings — and roughly how many users and items?

#167 **EXECUTABLE**

confidence 0.95

Add streaming to the chat endpoint in app/api/chat/route.ts — stream the Claude response as SSE tokens to the client instead of waiting for the full completion.

#168 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.65

reduce llm costs

Ambiguities

- Acceptable quality tradeoff unknown
- Current spend breakdown by call site is discoverable from logs, but quality bar isn't
- Options range from caching to model downgrades to prompt trimming

Missing information

- How much quality regression is acceptable and on which flows

Clarifying question: Are any of the LLM features allowed to get slightly worse to save money, or do all flows need current quality — in which case I'll focus on prompt caching and trimming rather than smaller models?

#169 **EXECUTABLE**

confidence 0.95

Our scikit-learn churn model in models/churn.pkl was trained on a dataframe with 14 features. Write the inference wrapper in serve/predict.py that validates incoming JSON has all features, applies the saved scaler from models/scaler.pkl, and returns the probability.

#170 **NEEDS CLARIFICATION** type: missing_context

confidence 0.80

the agent keeps going off the rails

Ambiguities

- Failure mode unspecified (infinite loops, wrong tools, ignoring instructions, unsafe actions)
- No transcripts provided
- Which agent/workflow is affected

Missing information

- Example transcripts of the bad behavior

Clarifying question: What does going off the rails look like — can you share a transcript or two? Looping, calling wrong tools, and ignoring instructions each get fixed differently.

#171 **EXECUTABLE**

confidence 0.94

Chunk the documents in data/docs/ for RAG: split on markdown headers first, then cap chunks at 800 tokens with 100 overlap, and store chunk metadata (source file, header path) alongside the text in the existing Chroma collection.

#172 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.90

we need to train on customer data without privacy issues

Ambiguities

- Applicable regime unspecified (GDPR, HIPAA, contractual DPAs)
- Whether anonymization, consent, or on-prem training is required
- What 'privacy issues' means to their legal/compliance team

Missing information

- The compliance constraints and data sensitivity classification

Clarifying question: What constraint are you operating under — GDPR/CCPA, customer contracts that forbid training on their data, or internal policy — because that decides between anonymization, opt-in consent, or not doing it at all?

#173 **EXECUTABLE**

confidence 0.95

Cache identical LLM requests in src/llm/client.py with a Redis-backed cache keyed on a hash of (model, messages, temperature), TTL 24h, with a bypass flag for callers that need fresh responses.

#174 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.86

label the dataset

Ambiguities

- Label taxonomy undefined
- Which dataset if multiple exist
- Whether LLM-assisted labeling is acceptable vs human-only
- Quality bar and ambiguous-case handling unspecified

Missing information

- The label set and labeling guidelines

Clarifying question: What are the possible labels and is there a guideline doc — or should I draft the taxonomy from a sample of the data first for you to approve?

#175 **EXECUTABLE**

confidence 0.94

Export the PyTorch model in models/classifier.pt to ONNX, verify the outputs match within 1e-4 on the test batch in tests/fixtures/batch.pt, and write the conversion script to scripts/export_onnx.py.

#176 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.85

make the data pipeline real time

Ambiguities

- 'Real time' latency target undefined (seconds? minutes?)
- Which outputs need freshness is unspecified
- Streaming infra (Kafka, Kinesis) may need introducing — big architectural change

Missing information

- Actual latency requirement and which data products need it

Clarifying question: What latency do consumers actually need — sub-second, or is a few minutes fine? Micro-batching every few minutes is a much smaller change than true streaming.

#177 **EXECUTABLE**

confidence 0.94

Add a system prompt A/B comparison mode to our eval script evals/run.py: take two prompt files, run both against the same cases, and output a side-by-side win/loss table judged by exact match.

#178 **NEEDS CLARIFICATION** type: missing_context

confidence 0.66

why are the predictions different every time

Ambiguities

- Which system — an LLM call (temperature) or an ML model (unseeded randomness)
- Whether determinism is actually required or just surprising
- Code is inspectable but the expectation isn't stated

Missing information

- Which prediction system and whether determinism is a requirement

Clarifying question: Is this the LLM-backed feature or the trained model — and do you need identical outputs run-to-run, or are you worried the variance indicates a bug?

#179 **EXECUTABLE**

confidence 0.95

Set up MLflow tracking for train.py: log params, per-epoch loss and accuracy, and the final model artifact to the local mlruns/ store.

#180 **NEEDS CLARIFICATION** type: missing_context

confidence 0.60

the gpu is sitting idle half the time during training

Ambiguities

- Likely a data-loading bottleneck but profiling evidence not provided — though the training code is inspectable
- Hardware setup and batch size unknown without looking
- Whether the user wants diagnosis or a fix

Missing information

- Profiler output or permission to instrument and rerun training

Clarifying question: Can I run a short profiled training run to confirm where the stall is — it's usually the data loader, but I want torch profiler output before changing num_workers, prefetching, or batch size?

#181 **EXECUTABLE**

confidence 0.96

Add pull-to-refresh to the feed screen in app/screens/FeedScreen.tsx using RefreshControl, calling the existing refetchFeed() from our React Query hook.

#182 **NEEDS CLARIFICATION** type: missing_context

confidence 0.85

the app crashes on android

Ambiguities

- No crash log or stack trace provided
- Which screen/action triggers it unknown
- Device/OS version pattern unknown

Missing information

- The crash log and reproduction steps

Clarifying question: Do you have the crash log from Logcat, Crashlytics, or the Play Console — or at least which screen or action triggers it?

#183 **EXECUTABLE**

confidence 0.92

Bump the Android targetSdkVersion to 35 in android/app/build.gradle, fix any manifest issues the upgrade flags, and verify the debug build compiles.

#184 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.95

make an app for our website

Ambiguities

- Native vs cross-platform vs PWA vs webview wrapper undecided
- Which site features the app needs
- Platforms (iOS, Android, both) unspecified
- App store distribution requirements unknown

Missing information

- Whether a wrapped webview suffices or native features are needed, and target platforms

Clarifying question: Does the app need native capabilities like push notifications, offline mode, or camera access — or is it mainly your website in app form, which could be a much simpler PWA or webview?

#185 **EXECUTABLE**

confidence 0.96

Add haptic feedback on the like button in ios/Sources/Views/PostCard.swift using UIImpactFeedbackGenerator with medium style.

#186 **NEEDS CLARIFICATION** type: missing_context

confidence 0.75

push notifications aren't working

Ambiguities

- Platform unspecified (iOS, Android, both)
- Failure point unknown (registration, sending, delivery, display)
- Environment matters (APNs sandbox vs prod certs)

Missing information

- Which platform and whether tokens register successfully

Clarifying question: Is this iOS, Android, or both — and do devices successfully register tokens (check your backend logs), or do tokens register but notifications never arrive?

#187 **EXECUTABLE**

confidence 0.92

Migrate the settings screen from UIKit to SwiftUI: rebuild `ios/Sources/Settings/SettingsViewController.swift` as a SwiftUI view with the same rows and behavior, hosted via `UIHostingController` so the rest of the app is untouched.

#188 **NEEDS CLARIFICATION** type: missing_context

confidence 0.66

the app is slow on older phones

Ambiguities

- Which screens/interactions are slow is unspecified
- How old is 'older' — minimum supported device unknown
- No profiling data, though profiling is possible with a target device/simulator

Missing information

- Which interactions are slow and the oldest device tier to support

Clarifying question: Which screens feel slow, and what's the oldest device you need to support well — that tells me what to profile against?

#189 **EXECUTABLE**

confidence 0.94

Add deep linking for `myapp://post/<id>` on both platforms in this React Native app — route it to the `PostDetail` screen via our existing React Navigation linking config in `app/navigation/linking.ts`.

#190 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

get us in the app store

Ambiguities

- App readiness unknown (signing, assets, privacy policy)
- Apple Developer account status unknown
- Which store(s) — App Store, Play, or both
- Release strategy (TestFlight first?) unspecified

Missing information

- Developer account status and what's already prepared (icons, screenshots, privacy policy)

Clarifying question: Do you already have an Apple Developer account and the store assets (icon, screenshots, privacy policy URL) — or am I starting the whole submission checklist from scratch?

#191 **EXECUTABLE**

confidence 0.96

Sanitize the HTML in user comments before rendering: we render `comment.body` with `dangerouslySetInnerHTML` in `src/components/Comment.tsx`. Add `DOMPurify` with an allowlist of basic formatting tags.

#192 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.60

is our app secure

Ambiguities

- Unbounded question — threat model undefined
- Whether this is for a compliance audit, a customer questionnaire, or general hygiene
- Depth expected (quick scan vs full review) unknown

Missing information

- What prompted the question and the threat model

Clarifying question: What's prompting this — a customer security questionnaire, an upcoming audit, or general concern? I can run a focused review either way, but the bar differs a lot.

#193 **EXECUTABLE**

confidence 0.94

Audit our Express routes for endpoints missing the `requireAuth` middleware — list every unauthenticated route in `src/routes/` and flag which ones expose user data.

#194 **NEEDS CLARIFICATION** type: missing_context

confidence 0.85

fix the vulnerability

Ambiguities

- Which vulnerability — no CVE, report, or scanner output referenced
- Severity and exploitability unknown

Missing information

- The vulnerability report or scanner finding

Clarifying question: Which vulnerability — can you share the scanner finding, CVE ID, or the report that flagged it?

#195 **EXECUTABLE**

confidence 0.90

npm audit shows a high-severity prototype pollution finding in a transitive dependency of `webpack-dev-server`. Check if we're actually exposed, and fix it with an override if upgrading isn't clean.

#196 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.70

add 2fa

Ambiguities

- Second factor unspecified (TOTP app, SMS, email codes, passkeys)
- Mandatory vs optional enrollment
- Recovery flow for lost devices undefined

Missing information

- Which second factor and whether it's required for all users

Clarifying question: Which second factor do you want — authenticator-app TOTP is the usual default over SMS — and is 2FA optional per user or enforced for everyone?

#197 **EXECUTABLE**

confidence 0.97

Move the hardcoded API keys in `config/settings.py` to environment variables, add them to `.env.example` with placeholder values, and update the README setup section.

#198 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

we failed the pentest

Ambiguities

- The pentest report findings are not provided
- Severity ranking and remediation deadlines unknown

Missing information

- The pentest report or at least the critical findings

Clarifying question: Can you share the pentest report — or paste the critical and high findings — so I can start remediating in severity order?

#199 **EXECUTABLE**

confidence 0.95

Add CSRF protection to the Flask app: enable Flask-WTF's CSRFProtect globally, add the token to the three forms in templates/, and exempt the /api/webhooks/stripe endpoint which is HMAC-verified.

#200 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.92

encrypt the data

Ambiguities

- Which data and where (at rest, in transit, application-level field encryption)
- Threat being defended against unspecified
- Key management approach undefined

Missing information

- Which data needs encryption and against what threat

Clarifying question: Which data, and what are you protecting against — a stolen database backup (at-rest/field encryption), network snooping (TLS), or admins seeing sensitive fields (application-level encryption)?

#201 **EXECUTABLE**

confidence 0.94

Set the security headers on our Next.js app via next.config.js: HSTS, X-Content-Type-Options, X-Frame-Options DENY, and a CSP that allows our own assets plus plausible.io for analytics.

#202 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.68

make sure users cant see each others stuff

Ambiguities

- Which resources need isolation is unspecified, though routes/queries can be audited
- Whether any cross-user sharing is legitimate (teams, orgs, admins)
- Whether a known leak exists or this is preventive

Missing information

- Whether there's a known leak and what sharing is intentional

Clarifying question: Did someone actually see another user's data — and is any cross-user access legitimate in your product, like team members or admins? Then I'll audit every query for missing ownership checks.

#203 **EXECUTABLE**

confidence 0.95

Replace MD5 password hashing in app/auth/passwords.py with bcrypt, and rehash each user's password transparently on their next successful login.

#204 **NEEDS CLARIFICATION** type: missing_context

confidence 0.70

someone is hammering our api

Ambiguities

- Attack pattern unknown (single IP, distributed, specific endpoint)
- Whether it's malicious or a misbehaving legitimate client
- Current mitigation options depend on infra (CDN, WAF, app-level)

Missing information

- Traffic pattern details: source IPs, endpoints, request rate

Clarifying question: Can you pull the access logs — is it one IP hitting one endpoint (block it now) or distributed traffic (needs rate limiting or WAF rules), and is your API behind Cloudflare or similar?

#205 **EXECUTABLE**

confidence 0.95

Run gitleaks over the full git history of this repo and give me a report of any committed secrets, grouped by file and rule.

#206 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.72

lock the admin panel down to just us

Ambiguities

- Who 'us' is — role-based, specific accounts, IP allowlist, VPN
- How admin access is currently controlled is inspectable, but the desired policy isn't stated

Missing information

- The access policy: which people/roles/networks should retain access

Clarifying question: Who exactly should keep access — specific user accounts, anyone with an admin role, or only requests from your office VPN/IP range?

#207 **EXECUTABLE**

confidence 0.97

Add input length limits and a request body size cap to the public contact form endpoint in api/contact.ts — max 200 chars name, 5000 chars message, reject bodies over 50KB with 413.

#208 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.94

gdpr compliance

Ambiguities

- GDPR spans data export, deletion, consent, retention, DPAs — which obligation is the ask
- Current state of compliance unknown
- Whether this is legal-driven with a deadline

Missing information

- Which GDPR obligations need implementing and the deadline

Clarifying question: Which part do you need first — user data export, right-to-deletion, cookie consent, or a data inventory/retention policy? If there's a legal letter or deadline driving this, that ranking changes.

#209 **EXECUTABLE**

confidence 0.96

Verify the Stripe webhook signatures in app/api/webhooks/stripe/route.ts using the signing secret from env, and reject requests older than 5 minutes.

#210 **NEEDS CLARIFICATION** type: missing_context

confidence 0.72

the login is broken but only for some users

Ambiguities

- What 'broken' looks like for affected users (error, wrong password message, infinite redirect)
- What the affected users share (SSO vs password, old accounts, specific email domain) is unknown
- Logs may reveal the pattern if accessible

Missing information

- The error affected users see and what they have in common

Clarifying question: What do the affected users see when login fails, and do they share anything — like all using Google SSO, or all having accounts older than some date?

#211 **EXECUTABLE**

confidence 0.93

Write an ADR in docs/adr/ documenting our decision to move from REST polling to WebSockets for order status updates — context, decision, alternatives considered (SSE, long polling), and consequences.

#212 **NEEDS CLARIFICATION** type: missing_context

confidence 0.75

should we use microservices

Ambiguities

- Current architecture and pain points unknown
- Team size and deployment maturity unstated
- What problem microservices would solve here is unclear

Missing information

- The specific pain the current architecture causes

Clarifying question: What's the pain today — deploys blocking each other, one component needing independent scaling, or team ownership conflicts? The answer usually points to a more targeted fix than full microservices.

#213 **EXECUTABLE**

confidence 0.91

Draw up a plan to split the billing module out of our Django monolith into its own service: identify the model boundaries from the code in billing/, list the cross-module imports that need API contracts, and propose a migration order. Don't change code yet.

#214 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.93

architect the system for scale

Ambiguities

- Scale target undefined (10x users? specific RPS?)
- Which dimension scales (traffic, data volume, team)
- Budget and timeline unknown

Missing information

- Concrete scale targets and which component hits limits first

Clarifying question: What scale are we designing for — roughly how many users/requests now versus in 12 months, and is the growth in traffic, data volume, or both?

#215 **EXECUTABLE**

confidence 0.93

Write the README for this repo: what it does (parse the CLI entry point and main modules to find out), setup steps from the Makefile, and how to run tests.

#216 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.96

i have an idea for an app like uber but for tutoring, build the mvp

Ambiguities

- Core MVP feature set undefined (matching, scheduling, payments, video?)
- Platform unspecified (web, mobile)
- What 'like Uber' means concretely (on-demand matching? ratings? geolocation?)

Missing information

- The minimal feature set for the first version

Clarifying question: For the very first version, what's the one core loop — students browse tutors and book a time slot, or on-demand instant matching — and is payment in or out of scope for the MVP?

#217 **EXECUTABLE**

confidence 0.90

Estimate the work to add SAML SSO to our app, broken down by task: we use Django with django-allauth, and the customer wants Okta. Give a task list with rough sizes, not code.

#218 **NEEDS CLARIFICATION** type: missing_context

confidence 0.93

how long will the migration take

Ambiguities

- Which migration is unspecified — no prior context available
- Scope and constraints of the migration unknown

Missing information

- What is being migrated from and to

Clarifying question: Which migration are you referring to — what's moving from where to where?

#219 **EXECUTABLE**

confidence 0.94

Write a postmortem doc for yesterday's outage using our template in docs/postmortems/TEMPLATE.md: the API was down 09:14-10:02 UTC because the cert renewal cron failed silently and the TLS cert expired. Timeline details are in the #incident-0612 export at docs/postmortems/incident-0612-log.txt.

#220 **NEEDS CLARIFICATION** type: missing_context

confidence 0.88

we keep having outages

Ambiguities

- Outage causes unknown — no incident history provided
- Whether causes repeat or vary determines the fix
- Current observability and on-call process unknown

Missing information

- What the last few outages were caused by

Clarifying question: What caused the last three outages — same root cause repeating, or different things each time? Same cause means fix that thing; varied causes means investing in observability and rollback speed.

#221 **EXECUTABLE**

confidence 0.93

Review PR #847 in our repo — run `gh pr checkout 847`, read the diff, and leave a review comment summary covering correctness and anything risky. Don't approve or merge.

#222 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.94

modernize the codebase

Ambiguities

- 'Modernize' is unbounded (language version, frameworks, patterns, tooling)
- Risk tolerance and timeline unknown
- What problem the old code causes is unstated

Missing information

- Which aspect of the codebase hurts most and the appetite for risk

Clarifying question: What hurts most right now — outdated framework versions blocking upgrades, untyped code causing bugs, or old patterns slowing the team — and how much breakage risk can you absorb at once?

#223 **EXECUTABLE**

confidence 0.92

Generate OpenAPI 3.1 docs for the routes in `src/routes/` – derive schemas from the zod validators already attached to each route and write the spec to `docs/openapi.yaml`.

#224 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

what stack should i use

Ambiguities

- What is being built is unknown
- Team experience and existing skills unstated
- Constraints (timeline, hosting budget, hiring) unknown

Missing information

- What the product is and the team's existing experience

Clarifying question: What are you building, and what languages does your team already know well — familiarity usually beats theoretical best choice?

#225 **EXECUTABLE**

confidence 0.95

Write a `CONTRIBUTING.md` covering: branch naming (`feature/*`, `fix/*`), conventional commits, PR review requirements (1 approval, CI green), and how to run the dev environment per the existing `docker-compose` setup.

#226 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.80

the new hire starts monday, get the codebase ready

Ambiguities

- 'Ready' is undefined — onboarding docs, dev env setup, good first issues, access?
- What the new hire will work on is unknown

Missing information

- What 'ready' means and what the new hire's first project is

Clarifying question: What will they work on first — and by 'ready' do you mean onboarding docs and a reliable dev-setup guide, a list of starter tasks, or both?

#227 **EXECUTABLE**

confidence 0.92

Compare our two candidate queue options for background jobs – we're a Python/Django shop on AWS – and write a one-page recommendation: Celery+SQS vs Celery+Redis, covering ops burden, retry semantics, and local dev experience.

#228 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.85

we need to be soc2 compliant by q3

Ambiguities

- Type I vs Type II not specified
- Which trust criteria beyond Security are in scope
- Current control maturity unknown
- Whether an auditor/platform (Vanta, Drata) is chosen

Missing information

- SOC 2 type, trust criteria in scope, and current tooling

Clarifying question: Is this SOC 2 Type I or Type II, and have you picked a compliance platform like Vanta or Drata — that determines whether I'm building evidence automation or starting with the control gap analysis?

#229 **EXECUTABLE**

confidence 0.92

Triage the open GitHub issues: label each of the ~40 open issues in this repo as bug/feature/question, flag duplicates, and list the five that look most critical with a one-line reason each.

#230 **NEEDS CLARIFICATION** type: missing_context

confidence 0.65

my cofounder says the code is bad, can you check

Ambiguities

- What the cofounder's specific complaints are is unknown
- 'Bad' criteria undefined — though a general code-quality review is possible
- Which areas to prioritize in a large codebase

Missing information

- The cofounder's specific concerns, if any were given

Clarifying question: Did your cofounder point at anything specific — bugs, slow features, unreadable modules — or should I do a general health review and report the top issues with severity?

#231 **EXECUTABLE**

confidence 0.93

Write a one-page tech spec for adding CSV export to the reports page: API endpoint, async generation for large exports with email delivery, and a size cutoff for sync vs async. Our stack is Rails + Sidekiq + S3.

#232 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.85

clean up the backlog

Ambiguities

- Which tracker and project is unspecified
- Cleanup rules undefined (close stale? merge dupes? reprioritize?)
- Authority to close others' tickets unclear

Missing information

- Where the backlog lives and what cleanup actions are authorized

Clarifying question: Where does the backlog live (Jira, Linear, GitHub), and am I allowed to close stale tickets outright or only flag them for your review?

#233 **EXECUTABLE**

confidence 0.95

Document the env vars: scan the codebase for every process.env and os.environ access, and produce docs/environment.md listing each variable, where it's used, and whether it's required or has a default.

#234 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.70

investor demo is tomorrow and search is embarrassing

Ambiguities

- What's embarrassing about search (relevance, speed, ugly UI, crashes)
- Whether a demo-path fix suffices vs a real fix
- Which queries will be demoed

Missing information

- What's wrong with search and which queries the demo will use

Clarifying question: What's the embarrassing part — bad results, slowness, or appearance — and do you know which searches you'll run in the demo, so I can make at least those bulletproof by tomorrow?

#235 **EXECUTABLE**

confidence 0.92

Map our service dependencies: parse the docker-compose files and the K8s manifests in infra/, and generate a Mermaid diagram of which services talk to which, saved to docs/architecture.md.

#236 **NEEDS CLARIFICATION** type: missing_context

confidence 0.92

we're acquiring a company, how do we merge the codebases

Ambiguities

- Both codebases' stacks and overlap unknown
- Whether 'merge' means integrate products, share infra, or literally combine repos
- Timeline and what must keep running unclear

Missing information

- What both stacks are and what 'merged' needs to mean for the business

Clarifying question: What does merged mean for the business — both products keep running separately on shared infra, one product absorbs the other's features, or one codebase gets retired?

#237 **EXECUTABLE**

confidence 0.94

Audit our open-source dependency licenses: list every production dependency with a license that isn't MIT/Apache-2.0/BSD, and flag any GPL/AGPL ones — we distribute a commercial binary.

#238 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.95

make us look bigger than we are

Ambiguities

- Unclear what artifact this refers to (website, docs, status page, API polish)
- What impression matters to whom (customers, investors, candidates)

Missing information

- What surface should convey maturity and to which audience

Clarifying question: Which surface are we polishing and for whom — the marketing site for customers, API docs for developer credibility, or something like a status page and SLA to look enterprise-ready?

#239 **EXECUTABLE**

confidence 0.95

Rewrite the error messages in src/cli/errors.ts to be actionable: each should say what went wrong, the likely cause, and the command or fix to try — keep the existing error codes.

#240 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.70

the consultant left and nobody understands the system

Ambiguities

- Which system and what's most urgent to understand (deploys? data flows? a specific module?)
- Whether docs, credentials, or knowledge is the gap
- What the team needs to do with it next

Missing information

- What the team needs to do with the system first

Clarifying question: What do you need to be able to do first — deploy it safely, fix a current bug, or change a feature? I'll reverse-engineer and document that path before mapping the whole system.

#241 **EXECUTABLE**

confidence 0.96

Add graceful shutdown to cmd/server/main.go: trap SIGTERM, stop accepting new connections, drain in-flight requests with a 30s timeout, then close the DB pool.

#242 **NEEDS CLARIFICATION** type: missing_file

confidence 0.82

port this to go

Ambiguities

- 'This' has no referent — no file or project specified
- Whether full behavior parity or a fresh idiomatic rewrite is wanted

Missing information

- What code to port

Clarifying question: Which code are we porting — a specific file, a module, or the whole service?

#243 **EXECUTABLE**

confidence 0.92

Fix the borrow checker error in src/parser.rs:142 — `cannot borrow self.tokens as mutable more than once` — restructure the loop so the second mutable borrow is scoped out.

#244 **NEEDS CLARIFICATION** type: missing_context

confidence 0.75

should i learn rust or go

Ambiguities

- Learner's goal unknown (job market, systems programming, web services, hobby)
- Current background unstated

Missing information

- What they want to build or what job they're aiming for

Clarifying question: What do you want to do with it — backend services and infrastructure tooling (Go gets you productive faster) or systems-level work like embedded, game engines, or performance-critical code (Rust)?

#245 **EXECUTABLE**

confidence 0.91

Upgrade this Spring Boot app from 2.7 to 3.2: update the parent pom, migrate javax.* imports to jakarta.*, fix the deprecated WebSecurityConfigurerAdapter usage in SecurityConfig.java, and get the build green.

#246 **NEEDS CLARIFICATION** type: missing_context

confidence 0.62

the java service eats all the cpu

Ambiguities

- When it happens (startup, steady load, specific requests) is unknown
- No thread dump or profiler data, though one can be captured if the service is reachable
- Could be GC thrashing, a spinning thread, or genuine load

Missing information

- When the CPU spikes occur and access to capture a thread dump

Clarifying question: Is the CPU pegged constantly or in bursts — and is the service running somewhere I can grab a thread dump and GC logs from?

#247 **EXECUTABLE**

confidence 0.94

Convert the callback-style code in lib/imageProcessor.js to async/await, keeping the public API backward compatible by still accepting an optional callback.

#248 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.66

upgrade php

Ambiguities

- Target version unspecified, though current version is discoverable
- Deployment environment controls what's installable
- Framework compatibility (e.g., old Laravel) may block the latest

Missing information

- Target PHP version and whether the hosting environment is under our control

Clarifying question: Upgrade to which version — latest 8.x — and do you control the server it runs on, or is the version pinned by your hosting provider?

#249 **EXECUTABLE**

confidence 0.97

Write a bash script scripts/backup_db.sh that pg_dumps \$DATABASE_URL to a timestamped gzip file in /var/backups/, keeps the 14 most recent, and exits nonzero with a message on any failure.

#250 **NEEDS CLARIFICATION** type: missing_context

confidence 0.96

automate the thing i do every morning

Ambiguities

- The morning routine is entirely unknown
- Tools/systems involved unspecified

Missing information

- What the manual routine consists of. step by step

Clarifying question: Walk me through the routine — what do you do step by step each morning, and in which tools?

#251 **EXECUTABLE**

confidence 0.94

Add structured logging to the Go service: replace the log.Printf calls in internal/ with slog, JSON handler, request ID from context, and levels – errors at Error, the rest at Info or Debug as appropriate.

#252 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.90

make a discord bot

Ambiguities

- Bot's purpose/commands undefined
- Language preference unknown
- Hosting plan unspecified

Missing information

- What the bot should do

Clarifying question: What should the bot do — moderate, respond to slash commands, post scheduled updates, or something else?

#253 **EXECUTABLE**

confidence 0.94

Profile scripts/generate_report.py with cProfile on the sample input in testdata/large.csv, and optimize the top hotspot — I suspect the row-by-row DataFrame appends.

#254 **NEEDS CLARIFICATION** type: missing_context

confidence 0.68

the script stopped working after the update

Ambiguities

- Which script and which update (OS, package, language runtime)
- Error output not provided, though running it may reproduce

Missing information

- Which script and what was updated

Clarifying question: Which script, and what got updated — a package, Python itself, the OS? If you point me at it I can run it and read the failure.

#255 **EXECUTABLE**

confidence 0.97

Write a PowerShell script that finds all files over 100MB under C:\Projects, lists them sorted by size, and exports the list to large_files.csv.

#256 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.78

do the regex thing for emails

Ambiguities

- Could mean validate emails, extract them from text, or mask/redact them
- Where it gets applied (which file/function) is unspecified
- Strictness of validation is a known rabbit hole

Missing information

- Whether this is validation, extraction, or redaction, and where

Clarifying question: Do you need to validate email input, extract emails from text, or redact them — and in which file?

#257 **EXECUTABLE**

confidence 0.95

Set up husky with a pre-push hook that runs npm test, and skip it when the PUSH_SKIP_TESTS env var is set.

#258 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

make a script to scrape the prices

Ambiguities

- Target site(s) unspecified
- Which products/pages to scrape unknown
- Output destination and frequency undefined
- Site's ToS/robots considerations depend on the target

Missing information

- Which site and which products to scrape

Clarifying question: Which site and which products — can you give me one example URL and tell me where the prices should end up (CSV, database, spreadsheet)?

#259 **EXECUTABLE**

confidence 0.93

Replace the recursive directory walk in src/indexer.cs with a parallel one using Parallel.ForEachAsync, capped at 8 concurrent workers, and keep the existing cancellation token support.

#260 **NEEDS CLARIFICATION** type: missing_file

confidence 0.80

turn this into a cli

Ambiguities

- 'This' has no referent
- Desired commands/flags undefined

Missing information

- Which code to wrap as a CLI

Clarifying question: Which code should become the CLI — point me at the file or function, and tell me what the main command should be called?

#261 **EXECUTABLE**

confidence 0.96

Add a retry decorator in utils/retry.py with configurable attempts, exponential backoff, and an exception allowlist; then apply it to the three flaky S3 calls in storage/s3.py.

#262 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.95

i need a chrome extension

Ambiguities

- Extension's purpose entirely unspecified
- Which sites it acts on unknown
- UI form (popup, content script, side panel) undefined

Missing information

- What the extension should do and on which sites

Clarifying question: What should the extension do, and on which sites — does it modify pages you visit, or add a tool you open from the toolbar?

#263 **EXECUTABLE**

confidence 0.93

Wrap the legacy SOAP client in services/legacy/soap_client.py with a typed facade: dataclasses for the three response types we use, timeouts, and translation of SOAP faults into our ApiError hierarchy.

#264 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

sync the data between the two systems

Ambiguities

- Which two systems is unspecified
- Direction of sync (one-way, two-way) undefined
- Conflict resolution and frequency unknown

Missing information

- Which systems, sync direction, and frequency

Clarifying question: Which two systems, and is it one-way (which is the source of truth?) or two-way — and how fresh does the synced data need to be?

#265 **EXECUTABLE**

confidence 0.94

Add tab completion for our CLI's subcommands and flags — it's a Python Click app in cli/main.py — and document the install step for bash and zsh in the README.

#266 **NEEDS CLARIFICATION** type: missing_context

confidence 0.91

hook it up to the spreadsheet

Ambiguities

- 'It' has no referent
- Which spreadsheet (Google Sheets, Excel file, where?)
- Direction of data flow undefined

Missing information

- What connects to which spreadsheet and which way data flows

Clarifying question: What's getting connected to which spreadsheet — is this app data being written out to a Google Sheet, or the sheet feeding data in?

#267 **EXECUTABLE**

confidence 0.94

Migrate this repo from npm to pnpm: convert the lockfile, update CI in .github/workflows/ci.yml to use pnpm with caching, and fix any scripts that assume npm.

#268 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.58

the build takes forever

Ambiguities

- Current and target build times unstated, though the build can be run and measured
- Which build (local dev, CI, production bundle)
- Acceptable tradeoffs (caching infra, splitting) unknown

Missing information

- Which build is the problem and what duration would be acceptable

Clarifying question: Which build hurts — local dev rebuilds or CI — and what would feel acceptable, so I know when to stop optimizing?

#269 **EXECUTABLE**

confidence 0.95

Vendor the single function we use from the abandoned left-pad-like package 'strutil' into src/lib/strutil.ts with its tests, remove the dependency, and update imports.

#270 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.55

do whatever you think is best with the error handling

Ambiguities

- Open-ended delegation but the target area is unspecified
- Whole codebase vs a specific module
- Existing conventions discoverable, but the boundary of the change isn't

Missing information

- Which part of the codebase this applies to

Clarifying question: Happy to use my judgment — but across the whole codebase or a specific area? A repo-wide error-handling overhaul is a big diff to review.

#271 **EXECUTABLE**

confidence 0.95

I've set up a new Postgres read replica. Update the Django config so reports/*.py queries route to the replica via a database router, while everything else stays on the primary. Replica connection string is in env as REPLICA_DATABASE_URL.

#272 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.72

We're building a notification system. Users should get notified when someone comments on their post, when they're mentioned, and when a post they commented on gets a new reply. Build it.

Ambiguities

- Delivery channel unspecified (in-app feed, email, push, all three)
- Real-time requirement unknown
- Batching/digest preferences undefined

Missing information

- Which delivery channels to implement

Clarifying question: Which channels do these notifications go to — an in-app notification feed, email, push, or some combination?

#273 **EXECUTABLE**

confidence 0.95

Implement the password reset flow: POST /api/auth/forgot sends a single-use token valid 1 hour via our existing Sendgrid mailer, GET link lands on /reset-password, POST /api/auth/reset validates token and sets the new password. Follow the patterns in src/auth/.

#274 **NEEDS CLARIFICATION** type: missing_component

confidence 0.60

Take the validation logic out of the controller and put it where it belongs.

Ambiguities

- Which controller is unspecified — repo may have many
- 'Where it belongs' presumes a convention that may or may not exist in this codebase

Missing information

- Which controller's validation to extract

Clarifying question: Which controller — and if the codebase doesn't already have a validation layer, should I create one (e.g., form objects or a validators module)?

#275 **EXECUTABLE**

confidence 0.96

Our webhook handler at POST /hooks/github sometimes processes the same delivery twice when GitHub retries. Add idempotency: store the X-GitHub-Delivery ID in the webhook_deliveries table and skip already-seen IDs.

#276 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.58

Build a CLI tool that takes a folder of markdown files and turns them into a website.

Ambiguities

- Reasonable defaults exist (static HTML, simple theme) but expectations are unstated
- Navigation/index behavior undefined
- Whether an existing SSG is acceptable vs writing from scratch

Missing information

- Whether wrapping an existing static site generator is acceptable

Clarifying question: Is this a learning/from-scratch project, or is the goal just the result — in which case I'd wrap an existing generator like Eleventy rather than reinventing one?

#277 **EXECUTABLE**

confidence 0.92

Hunt down why the /api/feed endpoint allocates 400MB per request — load the heap profile in profiles/feed.heap.pb.gz with pprof, identify the allocation site, and fix it.

#278 **NEEDS CLARIFICATION** type: missing_scope

confidence 0.62

Add the ability for users to export their data.

Ambiguities

- Which data counts as 'their data' (posts only? account info? activity logs?)
- Format unspecified (JSON, CSV, zip)
- Whether this is GDPR-driven (completeness bar is higher)

Missing information

- What data is included in an export and the driver behind the feature

Clarifying question: Is this a GDPR-style full export of everything we hold about the user, or a practical export of their content — and does the format matter to whoever asked for it?

#279 **EXECUTABLE**

confidence 0.94

Replace the polling in src/hooks/useJobStatus.ts (currently every 2s) with our existing WebSocket channel — subscribe to job:{id} events, keep polling as a fallback when the socket is disconnected.

#280 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

Set up the project like the last one.

Ambiguities

- 'The last one' is not visible from this session
- Which aspects to replicate (structure, tooling, CI, deps) is unspecified

Missing information

- Which previous project to mirror and access to it

Clarifying question: Which project should I mirror — can you give me its repo path or URL so I can copy its structure and tooling?

#281 **EXECUTABLE**

confidence 0.96

Three of our Celery tasks in tasks/email.py don't set retries. Add autoretry_for=(ConnectionError, TimeoutError), max 5 retries with exponential backoff and jitter, matching how tasks/billing.py does it.

#282 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.90

Move us off Heroku.

Ambiguities

- Destination platform unspecified
- Motivation unknown (cost, scale, features) — shapes the right target
- Downtime tolerance for cutover undefined

Missing information

- The target platform and what's driving the move

Clarifying question: What's driving the move and where to — cost (Render/Fly are the easy lift), or do you need raw AWS/GCP for something Heroku can't do?

#283 **EXECUTABLE**

confidence 0.93

The signup conversion funnel event 'signup_completed' stopped firing after the checkout refactor in PR #312. Trace where the analytics.track call got dropped and restore it with the same properties as before.

#284 **NEEDS CLARIFICATION** type: missing_component

confidence 0.88

Combine the two services.

Ambiguities

- Which two services is unspecified if the repo has several
- 'Combine' could mean merge codebases, share a deployment, or merge databases

Missing information

- Which services and what combined means operationally

Clarifying question: Which two services — and does combining mean one codebase and one deployment, or keeping the code separate but co-deploying them?

#285 **EXECUTABLE**

confidence 0.96

Add optimistic UI to the todo toggle in src/components/ToDoItem.tsx: flip the checkbox immediately, roll back with a toast if the PATCH fails. We use TanStack Query — use its onMutate/onError pattern.

#286 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.85

Make the import handle the new file format too.

Ambiguities

- The new format is not described or exemplified
- Which import pipeline if several exist
- Whether old format support continues

Missing information

- A sample file or spec of the new format

Clarifying question: Can you share a sample file of the new format — or point me to its spec — so I know the structure to parse?

#287 **EXECUTABLE**

confidence 0.94

Kill the feature flag enable_new_editor: it's been 100% rolled out for two months. Remove the flag checks, delete the old editor code paths in src/editor/legacy/, and clean up the flag from config.

#288 **NEEDS CLARIFICATION** type: missing_context

confidence 0.87

Use the better algorithm here.

Ambiguities

- 'Here' has no file/function referent
- 'Better' algorithm unidentified — better at what is unknown

Missing information

- Which code location and what dimension of better (speed, memory, accuracy)

Clarifying question: Which function are we talking about, and better how — faster, less memory, or more accurate results?

#289 **EXECUTABLE**

confidence 0.95

Instrument the checkout flow with OpenTelemetry spans: wrap cart validation, payment intent creation, and order persistence in services/checkout.py as child spans of the request span, with the order total and item count as attributes.

#290 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.95

Support teams.

Ambiguities

- Entire team model undefined (roles, invites, shared resources, billing)
- Could even mean Microsoft Teams integration rather than a teams feature
- Migration of existing single-user data unclear

Missing information

- What a team is in this product and what members can share

Clarifying question: To confirm — multi-user team workspaces in the product (not Microsoft Teams integration)? If so, what do team members share: all data, selected projects, and who can invite or bill?

#291 **EXECUTABLE**

confidence 0.96

Backfill slugs for the 12k existing articles: generate from title with the existing slugify() in utils/text.py, dedupe collisions with -2/-3 suffixes, write a reversible data migration, and don't touch articles that already have slugs.

#292 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

It should behave like before but faster.

Ambiguities

- 'It' and 'before' lack referents — no prior context in this session
- Performance target unstated

Missing information

- What component this refers to and what changed

Clarifying question: What's the 'it' — which feature or function, and what changed that made it different from before?

#293 **EXECUTABLE**

confidence 0.92

Our Terraform plan shows it wants to destroy and recreate the RDS instance because of the identifier change in the last commit. Figure out how to apply the rename without replacement — moved block or state mv — and verify with terraform plan showing no destroy.

#294 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.65

Add the standard stuff to the new endpoint.

Ambiguities

- 'Standard stuff' is team-specific — auth? validation? logging? rate limits?
- Which new endpoint, though git diff may show it
- Conventions may be inferable from sibling endpoints but the user's intent set isn't confirmed

Missing information

- What the team considers standard for endpoints

Clarifying question: By standard stuff do you mean what the sibling endpoints have — auth middleware, zod validation, and request logging — or is there more on your list?

#295 **EXECUTABLE**

confidence 0.93

Diff our package.json against what's actually imported in src/ and remove dependencies nothing uses; run the build and tests after to confirm.

#296 **NEEDS CLARIFICATION** type: missing_requirement

confidence 0.93

Make the report match what finance wants.

Ambiguities

- Finance's requirements are not stated or attached
- Which report is unspecified
- Format/columns/period expectations unknown

Missing information

- Finance's actual requirements for the report

Clarifying question: Do you have finance's spec or an example of what they expect — even a screenshot of the format they want — and which report is this?

#297 **EXECUTABLE**

confidence 0.95

Hide the beta features behind the existing FeatureGate component for any user whose plan is 'free' — the beta feature list is in config/features.ts under BETA_FEATURES.

#298 **NEEDS CLARIFICATION** type: missing_file

confidence 0.68

Speed up the dashboard query. It joins users, orders, and events, and finance runs it every morning. Don't break the numbers.

Ambiguities

- Where the query lives is unspecified (app code, BI tool, saved SQL?)
- 'Don't break the numbers' implies a correctness bar but no reference output to validate against

Missing information

- Where the query is defined and a reference result to validate against

Clarifying question: Where does that query live — in the app code, a BI tool, or a saved SQL file — and can you give me a known-correct output snapshot to validate my optimized version against?

#299 **EXECUTABLE**

confidence 0.95

Wire up Sentry in the React app: DSN from env, source maps uploaded in the CI build step, errors tagged with release version, and the ErrorBoundary in src/App.tsx reporting to it.

#300 **NEEDS CLARIFICATION** type: missing_context

confidence 0.90

Just do what the ticket says.

Ambiguities

- No ticket reference, link, or content provided
- Tracker system unknown

Missing information

- The ticket ID or its contents

Clarifying question: Which ticket — can you paste its contents or give me the ID and tracker so I can pull it up?

#301 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.72

send an sms to the user when their order ships

Ambiguities

- SMS provider unspecified (Twilio, AWS SNS, Vonage)
- No existing SMS integration may exist — account/credentials needed
- International sending requirements unknown

Missing information

- Which SMS provider to use and whether an account exists

Clarifying question: Do you already have an SMS provider account like Twilio, or should I set one up — and do you need international delivery?

#302 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.65

set up error tracking

Ambiguities

- Service unspecified (Sentry, Rollbar, Bugsnag, self-hosted)
- Budget for a paid tier unknown
- Which apps/services to instrument first

Missing information

- Preferred error-tracking service or constraints that pick one

Clarifying question: Any preference on the service — Sentry is the usual default and has a free tier — or a requirement like self-hosting that would rule SaaS out?

#303 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.60

the image processing should happen in a background queue

Ambiguities

- Queue backend unspecified (Redis+worker lib, SQS, RabbitMQ, DB-backed)
- Existing infra determines the cheap option — partially discoverable
- Delivery guarantees needed are unstated

Missing information

- What queue infrastructure is acceptable to introduce

Clarifying question: Is there already Redis or a cloud queue in the stack I should build on, or am I free to introduce one — and is it acceptable to lose a job on a crash, or do you need durable delivery?

#304 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.70

we need email – signup confirmations, password resets, that kind of thing

Ambiguities

- Email provider unspecified (SES, Sendgrid, Postmark, Resend)
- Sending domain and DNS access needed for deliverability
- Template/design expectations unknown

Missing information

- Email provider choice and access to DNS for domain verification

Clarifying question: Do you have an email provider account already, and can you change DNS records on your domain — deliverability setup (SPF/DKIM) needs that regardless of provider?

#305 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.66

add realtime updates to the dashboard so people don't have to refresh

Ambiguities

- Transport choice (WebSockets, SSE, short polling) depends on infra and scale
- Hosting platform may not support long-lived connections
- Freshness requirement unstated — seconds vs near-instant

Missing information

- Where the app is hosted and how fresh the data must be

Clarifying question: Where is the backend hosted — serverless platforms make WebSockets awkward, in which case SSE or a service like Pusher is the saner route — and is a few seconds of lag fine?

#306 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.68

we need a cms so marketing can edit the blog without us

Ambiguities

- CMS approach unspecified (headless like Sanity/Contentful, git-based, WordPress)
- Whether marketing needs a visual editor or markdown is fine
- Budget and hosting constraints unknown

Missing information

- Editor experience marketing expects and budget constraints

Clarifying question: Does marketing need a polished visual editor with previews (points to Sanity/Contentful), or are they comfortable with simple markdown forms (a git-based CMS is free and simpler)?

#307 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.70

transcribe the audio files users upload

Ambiguities

- Transcription approach unspecified (cloud API like Whisper API/Deepgram vs self-hosted model)
- Cost and privacy constraints unknown
- Latency requirement (realtime vs batch) unstated

Missing information

- Whether audio can leave your infrastructure and the latency requirement

Clarifying question: Can the audio be sent to a third-party API, or does privacy require on-prem transcription — and do users need results live or is processing within a few minutes fine?

#308 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.70

hook up analytics so we know what users actually do

Ambiguities

- Tool unspecified (PostHog, Mixpanel, Amplitude, GA4, self-hosted)
- Privacy/cookie-consent constraints unknown
- Which events matter is undefined

Missing information

- Analytics tool preference and the first questions it must answer

Clarifying question: What's the first question you want answered — funnel drop-off, feature usage, traffic sources — and do you have privacy constraints that rule out third-party trackers?

#309 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.75

the team needs a shared dev environment

Ambiguities

- Could mean devcontainers, a shared staging server, cloud dev envs (Codespaces), or just better docker-compose
- The actual pain being solved is unstated
- Budget unknown

Missing information

- What problem sharing solves — onboarding speed, parity, or shared test data

Clarifying question: What's the pain — 'works on my machine' drift (devcontainers fix that), or needing a shared running instance with common test data (that's a staging environment)?

#310 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.68

add maps to the store locator page

Ambiguities

- Map provider unspecified (Google Maps, Mapbox, Leaflet+OSM) — billing differs a lot
- API key/billing account availability unknown
- Features needed (geocoding, directions, clustering) unstated

Missing information

- Map provider preference and whether a billing account exists

Clarifying question: Do you have a Google Maps or Mapbox account with billing set up, or should I use the free Leaflet/OpenStreetMap route — and do you need directions/geocoding or just pins?

#311 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.72

we should cache at the edge

Ambiguities

- CDN/edge platform unspecified and current hosting unknown
- Which content is cacheable (static vs personalized) unanalyzed
- Invalidation requirements unstated

Missing information

- Current CDN/hosting setup and which responses are safe to cache

Clarifying question: Are you already behind a CDN like Cloudflare or CloudFront — and is the goal caching static assets, or full HTML/API responses (which raises cache-invalidation and personalization questions)?

#312 **NEEDS CLARIFICATION** type: missing_technology

confidence 0.70

enterprise customers are asking for sso

Ambiguities

- Protocol unspecified (SAML vs OIDC) — depends on customers' IdPs
- Build vs buy (WorkOS, Auth0) undecided
- Per-customer config (multi-tenant SSO) requirements unknown

Missing information

- Which IdPs customers use and build-vs-buy appetite

Clarifying question: Which identity providers do these customers use — Okta, Azure AD, Google Workspace — and are you open to a service like WorkOS, or does this need to be built in-house?

#313 **NEEDS CLARIFICATION** type: missing_file

confidence 0.85

add types to this

Ambiguities

- 'This' has no referent — no file or snippet attached
- Strictness expectations unknown

Missing information

- Which file or code to type

Clarifying question: Which file should I add types to?

#314 **NEEDS CLARIFICATION** type: missing_file

confidence 0.85

split this into smaller functions

Ambiguities

- 'This' has no referent in the session
- Decomposition preferences unknown

Missing information

- Which function or file to decompose

Clarifying question: Which function are we splitting — what file is it in?

#315 **NEEDS CLARIFICATION** type: missing_file

confidence 0.70

translate the comments to english

Ambiguities

- Which file(s) contain non-English comments is unspecified, though a scan could find them
- Whole repo vs specific module unclear

Missing information

- Which files to translate

Clarifying question: Should I scan the whole repo for non-English comments and translate everything I find, or is this about a specific file?

#316 **NEEDS CLARIFICATION** type: missing_file

confidence 0.83

the config is wrong, can you fix it

Ambiguities

- Which config file is unspecified — repos typically have many
- What's wrong with it is unstated

Missing information

- Which config file and what the correct value should be

Clarifying question: Which config file, and what's wrong — what behavior tells you it's misconfigured?

#317 **NEEDS CLARIFICATION** type: missing_file

confidence 0.80

rewrite it in typescript

Ambiguities

- 'It' has no referent — file, module, or whole project unclear
- Migration scope dramatically different per interpretation

Missing information

- What code to convert

Clarifying question: What's 'it' — a specific file, one package, or the entire JavaScript codebase?

#318 **NEEDS CLARIFICATION** type: missing_file

confidence 0.85

add docstrings to the file

Ambiguities

- 'The file' is not identified
- Docstring style (Google, NumPy, Sphinx) unspecified though repo convention may exist

Missing information

- Which file to document

Clarifying question: Which file needs the docstrings?

#319 **NEEDS CLARIFICATION** type: missing_file

confidence 0.78

there's a typo in the docs somewhere, my coworker mentioned it

Ambiguities

- Neither the typo nor its location is known
- Secondhand report with no specifics — can't grep for an unknown word

Missing information

- The misspelled word or the page it's on

Clarifying question: Did your coworker say which page or what the typo was — even the word — so I'm not proofreading the entire docs site?

#320 **NEEDS CLARIFICATION** type: missing_file

confidence 0.92

apply the changes we discussed

Ambiguities

- No record of the discussed changes exists in this session
- Files affected unknown

Missing information

- The list of changes that were discussed

Clarifying question: I don't have that discussion in this session — can you list the changes you want applied?

#321 **NEEDS CLARIFICATION** type: missing_file

confidence 0.62

merge the duplicate functions

Ambiguities

- Which duplicates the user has in mind is unspecified, though a similarity scan could find candidates
- Risk tolerance for behavioral merging of near-duplicates unknown

Missing information

- Which functions the user considers duplicates

Clarifying question: Do you have specific functions in mind, or should I scan for near-identical code and show you the candidates before merging anything?

#322 **NEEDS CLARIFICATION** type: missing_component

confidence 0.80

the second button doesn't do anything

Ambiguities

- 'Second button' on which screen is unidentified
- Expected behavior of the button unknown

Missing information

- Which screen/button and what it should do

Clarifying question: Which screen is the button on, and what should it do when clicked?

#323 **NEEDS CLARIFICATION** type: missing_component

confidence 0.70

bump the version

Ambiguities

- Which package (monorepo may have several) is unspecified
- Bump size (patch/minor/major) unstated
- Whether to tag/publish as part of the bump unclear

Missing information

- Which package and what bump level

Clarifying question: Which package, and is this a patch, minor, or major bump — and should I tag/publish or just change the version number?

#324 **NEEDS CLARIFICATION** type: missing_component

confidence 0.78

the table needs sorting

Ambiguities

- Which table in the UI is unspecified
- Sortable columns and default sort undefined
- Client-side vs server-side sorting matters for large data

Missing information

- Which table and which columns should sort

Clarifying question: Which table — and should all columns be sortable or specific ones, with what default order?

#325 **NEEDS CLARIFICATION** type: missing_component

confidence 0.80

disable the animation, it's annoying

Ambiguities

- Which animation is unidentified — apps typically have several
- Disable entirely vs make subtler is unclear

Missing information

- Which animation and where it appears

Clarifying question: Which animation — where do you see it and on what action does it trigger?

#326 **NEEDS CLARIFICATION** type: missing_component

confidence 0.75

remove the banner

Ambiguities

- Which banner (cookie notice, promo, announcement) is unspecified
- Remove permanently vs hide conditionally unclear

Missing information

- Which banner to remove

Clarifying question: Which banner — the cookie notice, the promo bar, or something else — and should it be gone entirely or just hidden for some users?

#327 **NEEDS CLARIFICATION** type: missing_component

confidence 0.72

the users endpoint is returning the wrong data

Ambiguities

- Multiple user-related endpoints may exist — which one is unclear
- 'Wrong' is not characterized (missing fields, stale data, wrong users)

Missing information

- The exact endpoint and an example of wrong vs expected output

Clarifying question: Which endpoint exactly — and can you show me one response that's wrong and what it should have contained?

#328 **NEEDS CLARIFICATION** type: missing_component

confidence 0.85

update the env var on the server

Ambiguities

- Which variable and its new value are unspecified
- Which server/environment is unidentified
- Restart/redeploy implications unknown

Missing information

- Variable name, new value, and target server

Clarifying question: Which variable, what's the new value, and on which server or environment?

#329 **NEEDS CLARIFICATION** type: missing_component

confidence 0.74

fix the spacing on the cards

Ambiguities

- Which card component/page is unspecified — several views may use cards
- What's wrong with the spacing (too tight, uneven, overflowing) is unstated

Missing information

- Which cards and what the spacing should look like

Clarifying question: Which page's cards, and what's off — too cramped, uneven gaps, or content touching the edges?

#330 **NEEDS CLARIFICATION** type: missing_component

confidence 0.82

swap the logo

Ambiguities

- New logo asset not provided
- Whether all logo placements (navbar, favicon, emails, social cards) should change is unclear

Missing information

- The new logo file and which placements to update

Clarifying question: Can you give me the new logo file — and should I update everywhere it appears (favicon, emails, social previews) or just the site header?

#331 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.80

roll back the deploy

Ambiguities

- Could mean redeploy the previous artifact, git revert the merge, or both
- If the deploy ran DB migrations, code rollback alone may break things
- Which environment is affected is unstated

Missing information

- Whether the bad deploy included migrations and which environment to roll back

Clarifying question: Did that deploy include database migrations — if so I need to handle those too, not just redeploy the previous version — and is this production?

#332 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.82

archive the project

Ambiguities

- Could mean archive the GitHub repo (read-only), export and delete, or archive a project entity inside an app
- Data retention expectations unknown

Missing information

- Which kind of archival is intended

Clarifying question: Do you mean mark the GitHub repo as archived/read-only, or actually export its data and decommission the infrastructure behind it?

#333 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.85

merge the branches

Ambiguities

- Which branches and merge direction unspecified
- Merge vs rebase vs squash preference unknown
- Conflict resolution authority unclear

Missing information

- Which branch merges into which

Clarifying question: Which branch into which — and if there are conflicts, should I resolve them myself or stop and show you?

#334 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.88

make the api public

Ambiguities

- Could mean expose it to the internet, open it to third-party developers with keys/docs, or remove auth entirely
- Security implications differ drastically between readings

Missing information

- What 'public' means here — reachable, documented for partners, or unauthenticated

Clarifying question: Public in what sense — reachable from the internet instead of internal-only, or opened up for third-party developers with API keys and docs? I want to make sure we're not removing authentication.

#335 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.82

duplicate the production environment

Ambiguities

- Could mean clone the infrastructure, copy the data, or both
- Copying prod data raises PII concerns
- Purpose (staging, DR, debugging) changes the right approach

Missing information

- Whether data comes along and what the copy is for

Clarifying question: Infrastructure only, or the data too — and if data, can it contain real user PII or does it need anonymizing for the copy's purpose?

#336 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.86

clean the database

Ambiguities

- Could mean delete test/junk data, drop unused tables, or maintenance like VACUUM
- Destructive readings need explicit confirmation
- Which environment is unstated

Missing information

- What kind of cleaning and on which environment

Clarifying question: What kind of clean — removing junk/test rows, dropping obsolete tables, or just maintenance like vacuum/reindex — and is this on production?

#337 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.78

restart the service

Ambiguities

- Which service and which environment is unspecified
- Could mean restart the process, redeploy, or reboot the host
- Production restarts may need coordination

Missing information

- Which service/environment and the reason for the restart

Clarifying question: Which service and where — and what's wrong, so I can confirm a restart is actually the fix rather than masking something we should look at first?

#338 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.85

share the codebase with the contractor

Ambiguities

- Could mean GitHub collaborator access, a zip export, or a scrubbed copy
- Secrets/history may need scrubbing first
- Access scope (read vs write, whole repo vs part) undefined

Missing information

- The sharing mechanism and access level intended

Clarifying question: Should the contractor get collaborator access to the repo (read or write?), or a one-time copy — and either way, want me to check the history for secrets first?

#339 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.75

freeze the dependencies

Ambiguities

- Could mean pin exact versions in the manifest, commit a lockfile, or vendor packages entirely
- Motivation (reproducibility, security, offline builds) picks the right one

Missing information

- What problem freezing should solve

Clarifying question: What's the goal — reproducible installs (a committed lockfile does it), or surviving a registry outage / supply-chain concerns (which points to vendoring)?

#340 **NEEDS CLARIFICATION** type: multiple_interpretations

confidence 0.93

take the site down

Ambiguities

- Could mean temporary maintenance mode or permanent decommission
- Irreversible if interpreted as teardown
- Timing and user messaging unspecified

Missing information

- Temporary or permanent, and what visitors should see

Clarifying question: Temporarily — a maintenance page while we do something — or permanently shutting it down? And if temporary, what should visitors see in the meantime?

#341 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.90

make onboarding smoother

Ambiguities

- No definition of smooth or evidence of current friction
- Whether this is user onboarding or developer onboarding is ambiguous too
- No metric to judge success

Missing information

- Where users currently drop off or complain

Clarifying question: Is this end-user onboarding in the product or developer onboarding to the codebase — and do you have drop-off data or specific complaints pointing at the friction?

#342 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.80

search results should be more relevant

Ambiguities

- No example queries with bad results provided
- Relevance is subjective without a reference standard
- No eval set to verify improvement

Missing information

- Example queries where results are bad and what should have ranked

Clarifying question: Can you give me three or four real queries with bad results and what should have come up first — that becomes the test set I tune against?

#343 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.75

improve the logging

Ambiguities

- What questions logs currently fail to answer is unknown
- Could mean more detail, more structure, less noise, or better tooling
- A recent painful debugging session may be the real driver

Missing information

- What you couldn't figure out from logs recently

Clarifying question: What happened recently that the logs couldn't answer — that incident tells me whether we need request tracing, more context per line, or less noise?

#344 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.82

make the api developer friendly

Ambiguities

- Friendliness undefined — docs, error messages, SDKs, consistency, sandbox?
- Who the developers are (internal, partners, public) shapes priorities
- No feedback from actual API consumers cited

Missing information

- What API consumers actually struggle with

Clarifying question: Have the developers using it said what hurts — bad docs, cryptic errors, inconsistent endpoints — or should I do a first-time-integration walkthrough myself and fix what I trip over?

#345 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.72

polish the ui before launch

Ambiguities

- 'Polish' is unbounded — could be spacing/alignment passes, loading states, empty states, copy
- Launch date / time budget unknown
- Which screens are in the launch path is unspecified

Missing information

- Time available and which screens matter for launch

Clarifying question: How much time do we have, and which screens will launch users actually see — I'll do a polish pass (alignment, loading/empty states, copy) on those first?

#346 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.70

the report should look professional

Ambiguities

- Which report and its current state unknown
- 'Professional' undefined — branding, layout, charts?
- Audience (exec, customer, regulator) shapes the bar

Missing information

- Which report, its audience, and an example of the desired look

Clarifying question: Who's the audience for this report, and do you have an example of one that looks right — even from another company — that I can match?

#347 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.68

make signup take less steps

Ambiguities

- Which steps are negotiable vs legally/functionally required is unknown
- Target step count unstated
- Whether deferred collection (ask later) is acceptable

Missing information

- Which currently-collected fields can be dropped or deferred

Clarifying question: Of what signup currently collects, what can we drop or ask for later — is anything there for legal/billing reasons that must stay upfront?

#348 **NEEDS CLARIFICATION** type: missing_acceptance_criteria

confidence 0.66

raise the code quality bar on the payments module

Ambiguities

- The bar is undefined — coverage %, type strictness, review rules, complexity limits?
- Whether this means fixing existing code or gating future changes
- Payments code is risk-sensitive — refactor appetite unknown

Missing information

- What measurable bar defines done and whether existing code gets refactored

Clarifying question: Does raising the bar mean refactoring the existing payments code now, or putting gates on future changes (coverage thresholds, stricter review) — touching working payment code has real risk?

#349 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.60

do something about the errors in the logs

Ambiguities

- Whether the goal is fixing root causes, reducing noise, or setting up alerting is unclear
- Which logs/services, though they may be inspectable
- Some logged errors may be expected behavior

Missing information

- Whether the errors are causing real problems or just noise

Clarifying question: Are these errors actually breaking things for users, or is the problem that they're noise drowning out real signals — fix the causes vs clean up the logging?

#350 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.85

customers are complaining, can you look into it

Ambiguities

- What customers are complaining about is unstated
- No tickets, messages, or reports referenced
- Could be bugs, performance, UX, or billing

Missing information

- What the complaints actually say

Clarifying question: What are they complaining about — can you paste a couple of the actual complaints or point me to the support tickets?

#351 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.78

we need to deal with tech debt

Ambiguities

- Which debt hurts is unspecified — the term covers everything
- Whether the deliverable is a prioritized inventory or actual fixes
- Time budget unknown

Missing information

- Whether to audit and prioritize, or fix something specific

Clarifying question: Is there a specific area where debt is actively slowing the team, or do you want me to audit the codebase and produce a prioritized hit list first?

#352 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.92

help me with the database

Ambiguities

- No task specified — could be schema design, a query, performance, migration, or an error
- Which database is unknown

Missing information

- What database task or problem needs help

Clarifying question: What's going on with the database — are you designing something new, debugging an error, or trying to speed something up?

#353 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.75

can you look at the payment code

Ambiguities

- 'Look at' has no goal — review for bugs, understand it, check a specific concern?
- A suspected problem may be the unstated motivation

Missing information

- What prompted the request and what the review should focus on

Clarifying question: What's prompting this — a suspected bug, an upcoming change you want me to scope, or a general review for issues?

#354 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.80

make the app better for power users

Ambiguities

- What power users need is unstated (shortcuts, bulk actions, API, density?)
- Who the power users are and what they do most is unknown

Missing information

- What workflows power users repeat that are currently slow

Clarifying question: What do your power users do over and over that's currently tedious — that tells me whether the win is keyboard shortcuts, bulk operations, or an API?

#355 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.85

the boss wants better numbers on the dashboard

Ambiguities

- 'Better numbers' could mean different metrics, more accurate data, or nicer presentation
- Which dashboard and what the boss actually asked for is secondhand and vague

Missing information

- What the boss actually said or expects to see

Clarifying question: Did the boss say what's wrong with the current numbers — are they missing metrics they care about, wrong/stale, or just hard to read?

#356 **NEEDS CLARIFICATION** type: unclear_goal

confidence 0.70

figure out the deploy situation

Ambiguities

- 'The situation' is undefined — could be broken deploys, missing process, or choosing a platform
- Desired outcome (diagnosis, fix, writeup) unclear

Missing information

- What about deploys is unsatisfactory

Clarifying question: What's wrong with deploys today — are they failing, manual and scary, or nonexistent — and do you want me to fix it or write up a recommendation first?

#357 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.62

compress the images users upload

Ambiguities

- Quality vs size tradeoff unspecified
- Whether originals must be kept is unknown
- Target format (WebP/AVIF conversion?) and max dimensions undefined

Missing information

- Acceptable quality loss and whether originals are retained

Clarifying question: How aggressive can compression be — is visible quality loss acceptable — and do we keep the originals or only store the compressed versions?

#358 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.88

add retries to the payment call

Ambiguities

- Retrying payments risks double-charging unless the call is idempotent
- Provider's idempotency-key support needs confirming
- Which failures are safe to retry (network timeout vs declined) is undefined

Missing information

- Whether the payment API call is idempotent or supports idempotency keys

Clarifying question: Does the payment provider call support idempotency keys — retrying a charge that actually succeeded would double-bill, so I need that guarantee before adding retries?

#359 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.72

limit api usage per customer

Ambiguities

- The actual limits are unspecified (requests/min? daily quota? per plan tier?)
- Behavior at the limit undefined (429, queue, throttle)
- Whether limits vary by pricing plan is unknown

Missing information

- The limit values and whether they differ by plan

Clarifying question: What are the limits — same for everyone or tiered by plan — and when a customer hits one, should they get a 429 immediately or be throttled?

#360 **NEEDS CLARIFICATION** type: missing_constraint

confidence 0.68

the export needs to handle big files

Ambiguities

- 'Big' is unquantified — 100MB and 50GB need different architectures
- Current failure mode at size unknown (timeout, OOM)
- Acceptable export latency for large files unstated

Missing information

- Realistic maximum size and how long an export may take

Clarifying question: How big is big — what's the largest realistic export — and is it okay if huge exports take minutes and arrive by email/download link instead of synchronously?

#361 **EXECUTABLE**

confidence 0.96

Add a confirmation dialog before account deletion in src/pages/Settings.tsx using our existing ConfirmModal component — require typing the word DELETE to enable the button.

#362 EXECUTABLE

confidence 0.96

Fix the off-by-one in the pagination: page 2 repeats the last item of page 1. The offset math is in `src/services/listings.py` line 88.

#363 EXECUTABLE

confidence 0.97

Write a systemd unit file for the worker: runs `/opt/app/venv/bin/python -m app.worker` as user `appuser`, restarts on failure with 5s delay, starts after `network-online.target`.

#364 EXECUTABLE

confidence 0.96

Add OpenGraph and Twitter card meta tags to the blog post template in `templates/post.html` using the post's title, excerpt, and `cover_image` fields.

#365 EXECUTABLE

confidence 0.95

Our Redis cache keys have no TTL and memory is creeping up. Add a 1-hour TTL to the product cache writes in `services/catalog.py` and a 10-minute TTL to the session-preview keys in `services/preview.py`.

#366 EXECUTABLE

confidence 0.94

Convert the synchronous file upload in `handlers/upload.go` to stream to S3 with multipart upload instead of buffering the whole file in memory.

#367 EXECUTABLE

confidence 0.95

Add a GET `/api/v1/exchange-rates` endpoint that proxies the ECB daily rates feed, caches it for 6 hours in memory, and returns rates relative to a `?base=` currency param defaulting to EUR.

#368 EXECUTABLE

confidence 0.95

There's a race in the singleton init in `src/lib/db.ts` – two concurrent requests both create a pool. Make initialization safe with a module-level promise.

#369 EXECUTABLE

confidence 0.96

Add a Storybook story for the Badge component in `src/components/ui/Badge.tsx` covering all four variants and the disabled state.

#370 EXECUTABLE

confidence 0.96

Write an alembic migration that renames the column `users.fullname` to `users.full_name` and update every reference in the SQLAlchemy models and queries.

#371 EXECUTABLE

confidence 0.95

Hook the contact form up to our existing POST `/api/contact` endpoint – it currently just `console.logs`. Show a success message on 200 and field errors from the 422 response body.

#372 EXECUTABLE

confidence 0.95

Replace the hand-rolled CSV parser in `import/parse.js` with `papaparse`, keeping the same output shape – the current one breaks on quoted fields containing commas.

#373 EXECUTABLE

confidence 0.97

Add log rotation for `/var/log/app/*.log` via logrotate: daily, keep 14 days, compress, copytruncate.

#374 EXECUTABLE

confidence 0.95

Implement the dead-letter handling TODO in `workers/consumer.py`: after 3 failed processing attempts, publish the message to the existing `orders.dlq` queue with the error attached, and ack the original.

#375 EXECUTABLE

confidence 0.95

Make the Dockerfile multi-stage: build stage with the full Go toolchain, final stage FROM scratch with just the binary and CA certs. Should shrink the image from ~900MB.

#376 EXECUTABLE

confidence 0.97

Add an empty state to the projects list in `src/pages/Projects.tsx`: when there are zero projects, show an illustration-free card with 'No projects yet' and a button that opens the existing `NewProjectModal`.

#377 EXECUTABLE

confidence 0.93

Fix the timezone bug: scheduled posts publish at the wrong hour because `schedule_time` is stored naive and compared against `UTC now()` in `jobs/publisher.py`. Store UTC and convert at the edges.

#378 EXECUTABLE

confidence 0.93

Add ``npm run typecheck`` running `tsc --noEmit`, fix the 23 existing type errors it surfaces, and add the script to CI after lint.

#379 EXECUTABLE

confidence 0.96

Extract the magic numbers in `src/game/physics.ts` (gravity, friction, max velocity) into a named constants object with doc comments explaining the units.

#380 EXECUTABLE

confidence 0.96

Write a script that checks every link in `docs/**/*.md` and reports dead ones – internal anchors and external URLs both, with a non-zero exit for CI use.

#381 EXECUTABLE

confidence 0.93

Our API returns dates in three different formats across endpoints. Standardize all serializers in `api/serializers/` to ISO 8601 UTC with trailing Z, and list which endpoints changed so we can warn the mobile team.

#382 EXECUTABLE

confidence 0.95

Add request ID propagation: generate a UUID in the ingress middleware if `X-Request-ID` is absent, attach it to the context, include it in every log line and in error responses.

#383 EXECUTABLE

confidence 0.94

Memcached is being retired company-wide. Swap the cache adapter in `lib/cache.rb` from Dalli to our Redis cluster (`REDIS_URL` env), keeping the same `get/set/delete` interface.

#384 EXECUTABLE

confidence 0.94

Build a /debug/queues admin page (behind the existing admin auth) showing each Sidekiq queue's depth, oldest job age, and failure count, auto-refreshing every 10s.

#385 EXECUTABLE

confidence 0.96

The webhook retry loop in src/webhooks/sender.ts retries forever on 4xx. Stop retrying on any 4xx except 408 and 429, and log the permanent failures to the webhook_failures table.

#386 EXECUTABLE

confidence 0.95

Add shellcheck to CI for the scripts/ directory and fix everything it flags at warning level or above.

#387 EXECUTABLE

confidence 0.94

Implement RFC 7807 problem+json error responses in the Spring Boot API: a @ControllerAdvice translating our exception hierarchy to type/title/status/detail, keeping existing status codes.

#388 EXECUTABLE

confidence 0.93

Migrate the three cron entries in the README from the old server to GitHub Actions scheduled workflows, preserving their schedules and pointing them at the existing make targets.

#389 EXECUTABLE

confidence 0.97

Add prefers-reduced-motion support: wrap the animations in src/styles/animations.css in the appropriate media query so they're disabled for users who opt out.

#390 EXECUTABLE

confidence 0.94

Split the 2000-line src/utils.py by topic into utils/strings.py, utils/dates.py, utils/io.py etc., keep backward-compatible imports via utils/__init__.py re-exports, and make sure tests pass.

#391 EXECUTABLE

confidence 0.92

run the linter and fix what it finds

#392 EXECUTABLE

confidence 0.96

whats in this repo? give me a quick tour

#393 EXECUTABLE

confidence 0.98

explain what this regex does: `^(?=[A-Z])(?=[A-Za-z\d@$!%*?&]){8,}$`

#394 EXECUTABLE

confidence 0.93

commit what's staged with a sensible message

#395 EXECUTABLE

confidence 0.96

why does git say 'detached HEAD' and how do i get back to my branch without losing the commit i just made

#396 EXECUTABLE

confidence 0.98

im new to python – whats the difference between a list and a tuple and when would i use each

#397 EXECUTABLE

confidence 0.98

convert this curl command to a python requests call: curl -X POST
https://api.example.com/v1/items -H 'Authorization: Bearer \$TOKEN' -d '{"name":"test"}'

#398 EXECUTABLE

confidence 0.98

find every TODO and FIXME in src/ and list them with file and line

#399 EXECUTABLE

confidence 0.93

the dev server won't start, it says EADDRINUSE :3000 – sort it out

#400 EXECUTABLE

confidence 0.94

write a .editorconfig matching how this codebase is already formatted

#401 EXECUTABLE

confidence 0.97

what's the biggest file in this repo and why is it so large?

#402 EXECUTABLE

confidence 0.97

Add a CODEOWNERS file: platform team owns infra/ and .github/, payments team owns src/billing/, everything else falls to @org/eng-leads.

#403 EXECUTABLE

confidence 0.95

diff develop against main and summarize what would ship if we cut a release right now

#404 EXECUTABLE

confidence 0.92

Take the inline SQL in get_revenue_report() in reports/queries.py and add an EXPLAIN ANALYZE-based comment documenting its current plan and cost against the dev database.

#405 EXECUTABLE

confidence 0.94

upgrade all the patch-level dependency updates (no minors or majors), run the tests, and tell me if anything broke

#406 EXECUTABLE

confidence 0.96

Add a Slack notification step to the deploy workflow that posts to the webhook in the SLACK_DEPLOY_WEBHOOK secret on failure only, with the run URL and commit message.

#407 EXECUTABLE

confidence 0.95

Make the CLI exit codes meaningful: 0 success, 1 user error, 2 internal error – right now everything exits 1. Entry point is cli/main.go.

#408 EXECUTABLE

confidence 0.95

There's a TOCTOU bug in save_upload() in storage/files.py – it checks existence then writes. Use atomic write-to-temp-and-rename instead.

#409 EXECUTABLE

confidence 0.95

Generate TypeScript types from the JSON schema files in schemas/*.json using json-schema-to-typescript, output to src/types/generated/, and add the generation step to the build.

#410 EXECUTABLE

confidence 0.94

the favicon is still the framework default, replace it with public/logo.svg properly (all the sizes/manifest bits)

#411 EXECUTABLE

confidence 0.95

Wrap the flaky third-party geocoding call in services/geo.py with a circuit breaker: open after 5 consecutive failures, half-open after 60s, and fall back to returning None so callers degrade gracefully.

#412 EXECUTABLE

confidence 0.96

Add ARIA labels and roles to the icon-only buttons in src/components/Toolbar.tsx – screen readers currently announce them as just 'button'.

#413 EXECUTABLE

confidence 0.93

Profile the bundle: run the analyzer, tell me the three biggest contributors to the main chunk, and code-split any route-level ones.

#414 EXECUTABLE

confidence 0.94

De-flake test_payment_timeout in tests/integration/test_payments.py – it depends on a real 2s sleep; replace with freezegun time control.

#415 EXECUTABLE

confidence 0.95

Read through migrations/ and tell me if any migration is irreversible – flag the ones without a down() and the ones whose down() would lose data.

#416 EXECUTABLE

confidence 0.95

Hash the user emails in the analytics events before they leave our backend – events are assembled in analytics/tracker.py; use SHA-256 with the existing ANALYTICS_SALT.

#417 EXECUTABLE

confidence 0.96

Add a docker-compose healthcheck to the postgres service and make the api service wait for it via depends_on condition service_healthy – right now the api races the db on startup.

#418 EXECUTABLE

confidence 0.94

Our Express app trusts X-Forwarded-For blindly for rate limiting. Set trust proxy correctly for one proxy hop (we're behind a single nginx) so clients can't spoof their IP.

#419 EXECUTABLE

confidence 0.93

Vendor the OpenAPI spec of the partner API (<https://partner.example.com/openapi.json>) into specs/, generate a typed client with openapi-typescript, and replace the hand-written fetch calls in src/partner/.

#420 EXECUTABLE

confidence 0.91

check if we're affected by the latest lodash CVE and patch if so